

# CS 310 : Scalable Software Architectures

*Class session on Tuesday, October 1<sup>st</sup>*



## October 2024

Sunday	Monday	Tuesday	Wednesday	Thursday	Friday	Saturday
		1	2	3	4	5
6	7	8	9	10	11	12
13	14	15	16	17	18	19
20	21	22	23	24	25	26
27	28	29	30	31		

www.a-printable-calendar.com

## Notes:


- *Focus this week:*
  - *Relational databases*
- *Class sessions **\*are\*** being recorded this week*
  - *Will be available under Panopto on Canvas*
- *Project 01 due Wednesday Oct 9<sup>th</sup> @ 11:59pm*
  - *Build a simple photo app using AWS*
  - *Part 01 has been released, configuration of AWS*
  - *Part 02 to be released, client-side programming*
- *Office hours started Monday*
- *Optional SQL homework will be posted*



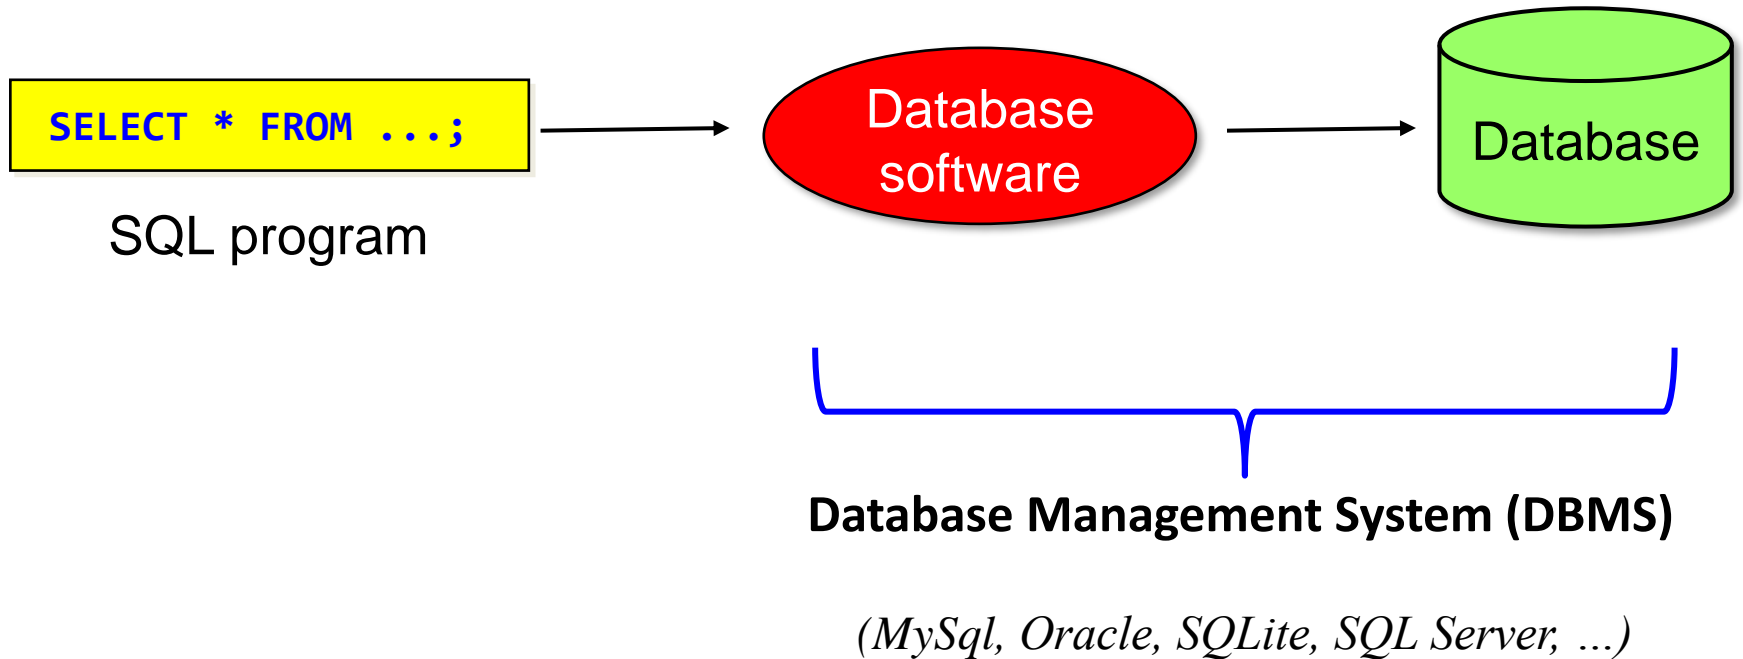
Northwestern  
University

# Goals for today

- **Summarize key points of lecture**
- **Write SQL**
  - *Using sqlite3, a local file-based DBMS (e.g. used in mobile apps)*
  - *Using MySQL running in AWS*
- **Work with Docker**
  - *You need Docker Desktop installed*
  - *Download files from GitHub:*
    - <https://github.com/joe-hummel/intro-to-db-sql>
    - Clone repo or download ZIP



# SQL + DBMS




SQLite is free, local , file based database, useful for local small devices like mobile.

# Select

- For retrieving data from a database
- General format:

```
SELECT <<the data you want>>
FROM   <<table>>
[ JOIN <<table> ON <<condition>> ]
[ WHERE      <<condition(s)>> ]
[ GROUP BY   <<one or more fields>> ]
[ HAVING     <<condition(s)>> ]
[ ORDER BY   <<one or more fields>> ]
;
```

optional



The **result** is  
\*always\* a  
table


# Group by

## Group By Ride\_Date

Sum(Num\_Riders)

Ride\_Date

41280	12/22/2017	W	6104
41000	12/18/2017	W	3636
40280	12/02/2017	A	1270
40140	12/18/2017	W	1759
40690	12/03/2017	U	499
41660	12/03/2017	A	8615
40180	12/03/2017	U	442
40250	12/22/2017	A	1353
40120	12/07/2017	W	3353
41420	12/22/2017	W	6034
40270	12/18/2017	A	887
41450	12/18/2017	W	9639
41210	12/02/2017	W	3210
40010	12/22/2017	U	641
41160	12/22/2017	U	621
40720	12/18/2017	W	613



41280	12/22/2017	W	6104
40250	12/22/2017	A	1353
41420	12/22/2017	W	6034
40010	12/22/2017	U	641
41160	12/22/2017	U	621

**14753**

41000	12/18/2017	W	3636
40140	12/18/2017	W	1759
40270	12/18/2017	A	887
41450	12/18/2017	W	9639
40720	12/18/2017	W	613

**16534**

40280	12/02/2017	A	1270
41210	12/02/2017	W	3210

**4480**

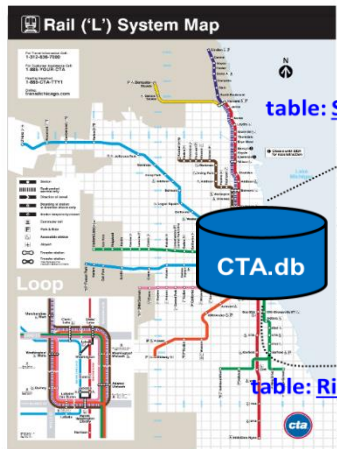
40690	12/03/2017	U	499
41660	12/03/2017	A	8615
40180	12/03/2017	U	442

**9556**

40120	12/07/2017	W	3353
-------	------------	---	------


**3353**

# Join



Station_ID	Station_Name
40010	Austin-Forest Park
40020	Harlem-Lake
40030	Pulaski-Lake
...	...

Station_ID	Ride_Date	Type_of_Day	Num_Riders
41280	2017-12-22 00:00:00.000	W	6104
40010	2017-12-28 00:00:00.000	W	1155
40280	2017-12-02 00:00:00.000	A	1270
40030	2017-12-24 00:00:00.000	U	595
...	...	...	...



```

** Top-10 Busiest Stations **
Lake/State|100,419,088
Clark/Lake|100,088,085
Chicago/State|91,899,932
Belmont-North Main|74,452,064
95th/Dan Ryan|74,235,360
Fullerton|72,888,906
Grand/State|68,379,115
O'Hare Airport|66,363,838
Jackson/State|61,803,911
Roosevelt|61,487,262

```

```
select "** Top-10 Busiest Stations **";
```

```


Select Station_Name, Sum(Num_Riders)
From Stations
Join Ridership On Stations.Station_ID = Ridership.Station_ID
Group By Stations.Station_ID
Order By Sum(Num_Riders) DESC
Limit 10;

```

# Getting the necessary software

## 1. Download files you need for today

- <https://github.com/joe-hummel/intro-to-db-sql>
- Clone repo or download ZIP



## 2. Make sure Docker Desktop is running

## 3. Build Docker image and run container:

### Linux/Mac/Windows WSL:

- 1) Open terminal, navigate to folder
- 2) `chmod 755 *.bash`
- 3) `./docker-build.bash`
- 4) `./docker-run.bash`

### Windows:

- 1) Open Powershell, navigate to folder
- 2) `.\docker-build.bat`
- 3) `.\docker-run.bat`

```
hummel> ./docker-run.bash
docker> ls
Dockerfile      datatier.py      docker-build.bat  docker-run.bat   main.sql
cta.db          docker-build.bash  docker-run.bash   main.py          movielens.db
docker>
```

# Common docker errors

## 1. "docker" command not found

- *Uninstall and reinstall Docker Desktop*

## 2. When you try to build, you are not authorized

- *docker login -u docker-username*

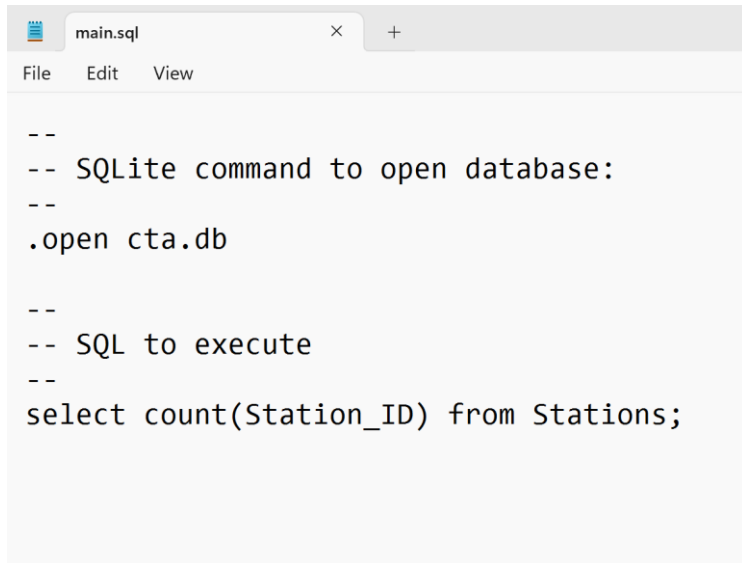
## 3. When you try to run, you get errors like "bash: \$\r: command not found"

1. *If you see the **docker**> prompt, type **exit***
2. `((Get-Content .bashrc) -join "`n") + "`n" | Set-Content -NoNewLine .bashrc`

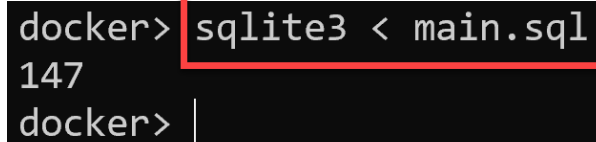


# Working with sqlite3

1. Open "main.sql" in a text editor
2. Write query and save changes
3. Run via docker container:

A screenshot of a text editor window titled 'main.sql'. The window has a menu bar with 'File', 'Edit', and 'View'. The content of the file is as follows:

```
--  
-- SQLite command to open database:  
--  
.open cta.db  
  
--  
-- SQL to execute  
--  
select count(Station_ID) from Stations;
```

A terminal window showing the execution of the sqlite3 command. The command 'sqlite3 < main.sql' is highlighted with a red box. The output '147' is shown on the next line.

```
docker> sqlite3 < main.sql  
147  
docker> |
```

# Exercise

- What is the yearly ridership for the "Noyes" station?

- Most DBMS have Year( ) function, sqlite does not

```
SELECT Station_Name,  
       strftime('%Y', Ride_Date) as Year,  
       Sum(Num_Riders)  
  
FROM   ? Ridership  
  
INNER JOIN ? Station on Ridership.Station_ID = Station.Station_ID  
  
WHERE  ? Station_Name = 'Noyes'  
  
GROUP BY ? Year  
  
ORDER BY ? Year ASC  
  
;
```

```
docker> sqlite3 < main.sql  
Noyes|2016|277442  
Noyes|2017|282461  
Noyes|2018|282356  
Noyes|2019|276037  
Noyes|2020|85834  
Noyes|2021|107387  
docker>
```

## Stations

Station_ID	Station_Name
40010	Austin-Forest Park
40020	Harlem-Lake
40030	Pulaski-Lake
...	...

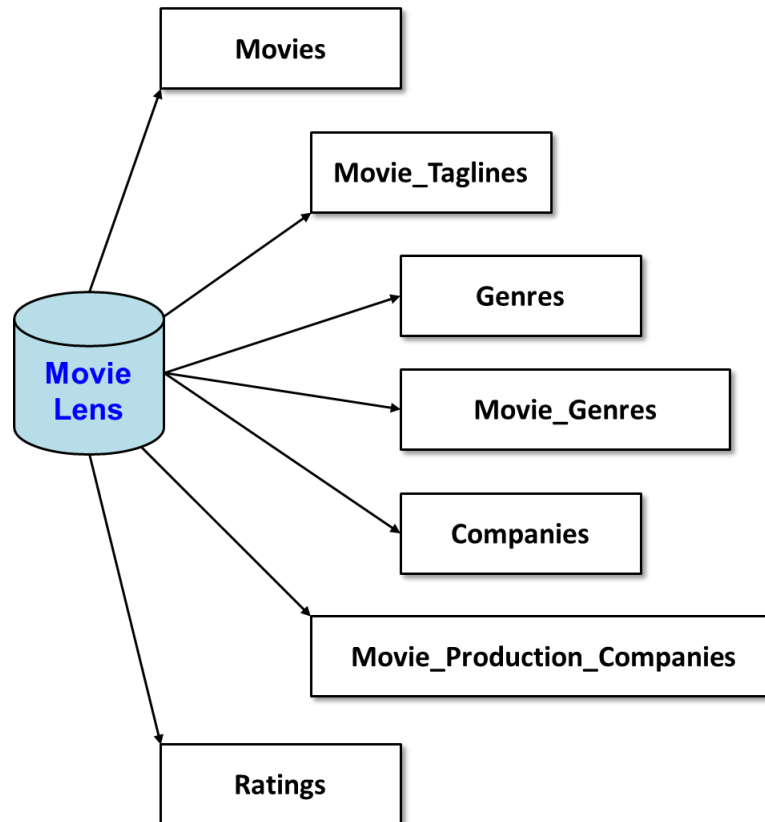
Station_ID	Ride_Date	Type_of_Day	Num_Riders
41280	2017-12-22 00:00:00.000	W	6104
40010	2017-12-28 00:00:00.000	W	1155
40280	2017-12-02 00:00:00.000	A	1270
40030	2017-12-24 00:00:00.000	U	595
...	...	...	...

## Ridership

# MovieLens database

- **MovieLens**

- <https://movielens.org/>



# Exercises

- How many movies are there? [45,431]
- How many ratings are there? [100,004]
- What is the movie id for 'The Matrix'? [ 603 ]
- Which movie titles contain 'matrix'? [ there are 7 ]

SELECT \* from Movies where title like '%matrix%';

## Movies

Movie_ID	Title	Release_Date	Runtime	Original_L anguage	Budget	Revenue
603	The Matrix	1999-03-30 00:00:00.000	136	en	63000000	463517383
862	Toy Story	1995-10-30 00:00:00.000	81	en	30000000	373554033

## Ratings

Movie_ID	Rating
605	8
603	6
605	10
605	6

# Exercise

```
--  
-- Retrieve the top-10 movies ranked by average rating;  
-- retrieve the title and average rating. Consider only  
-- movies with more than 100 reviews.  
--
```

```
SELECT ?  
FROM ?  
INNER JOIN ?  
GROUP BY ?  
HAVING ?  
ORDER BY ?  
LIMIT 10;
```

```
Sleepless in Seattle|8.975  
The Million Dollar Hotel|8.97427652733119  
Once Were Warriors|8.60655737704918  
Men in Black II|8.51339285714286  
Terminator 3: Rise of the Machines|8.51234567901234  
Confession of a Child of the Century|8.47107438016529  
The Thomas Crown Affair|8.47008547008547  
Shriek If You Know What I Did Last Friday the Thirteenth|8.454  
Scarface|8.44915254237288  
The 39 Steps|8.44329896907217
```

Movies

Movie_ID	Title	Release_Date	Runtime	Original_L anguage	Budget	Revenue
603	The Matrix	1999-03-30 00:00:00.000	136	en	63000000	463517383
862	Toy Story	1995-10-30 00:00:00.000	81	en	30000000	373554033

Ratings

Movie_ID	Rating
605	8
603	6
605	10
605	6

# Solution

```
--  
-- Retrieve the top-10 movies ranked by average rating;  
-- retrieve the title and average rating. Consider only  
-- movies with more than 100 reviews.  
--
```

```
SELECT Title, avg(Rating) as AvgRating  
FROM Movies  
INNER JOIN Ratings ON Movies.Movie_ID = Ratings.Movie_ID  
GROUP BY Ratings.Movie_ID  
HAVING Count(Rating) > 100  
ORDER BY AvgRating DESC, Title ASC  
LIMIT 10;
```

```
Sleepless in Seattle|8.975  
The Million Dollar Hotel|8.97427652733119  
Once Were Warriors|8.60655737704918  
Men in Black II|8.51339285714286  
Terminator 3: Rise of the Machines|8.51234567901234  
Confession of a Child of the Century|8.47107438016529  
The Thomas Crown Affair|8.47008547008547  
Shriek If You Know What I Did Last Friday the Thirteenth|8.454  
Scarface|8.44915254237288  
The 39 Steps|8.44329896907217
```

# Demo: My SQL Workbench

- In the projects we're going to work with MySQL running in AWS
  - *Need different software to connect to DB server*

Setup New Connection

Connection Name: MySQL running in AWS

Connection Method: Standard (TCP/IP)

Parameters SSL Advanced

Hostname: nu-cs-mysql.cb1xaky37wq8.us-east-2.rds.amazonaws.com Port: 3306

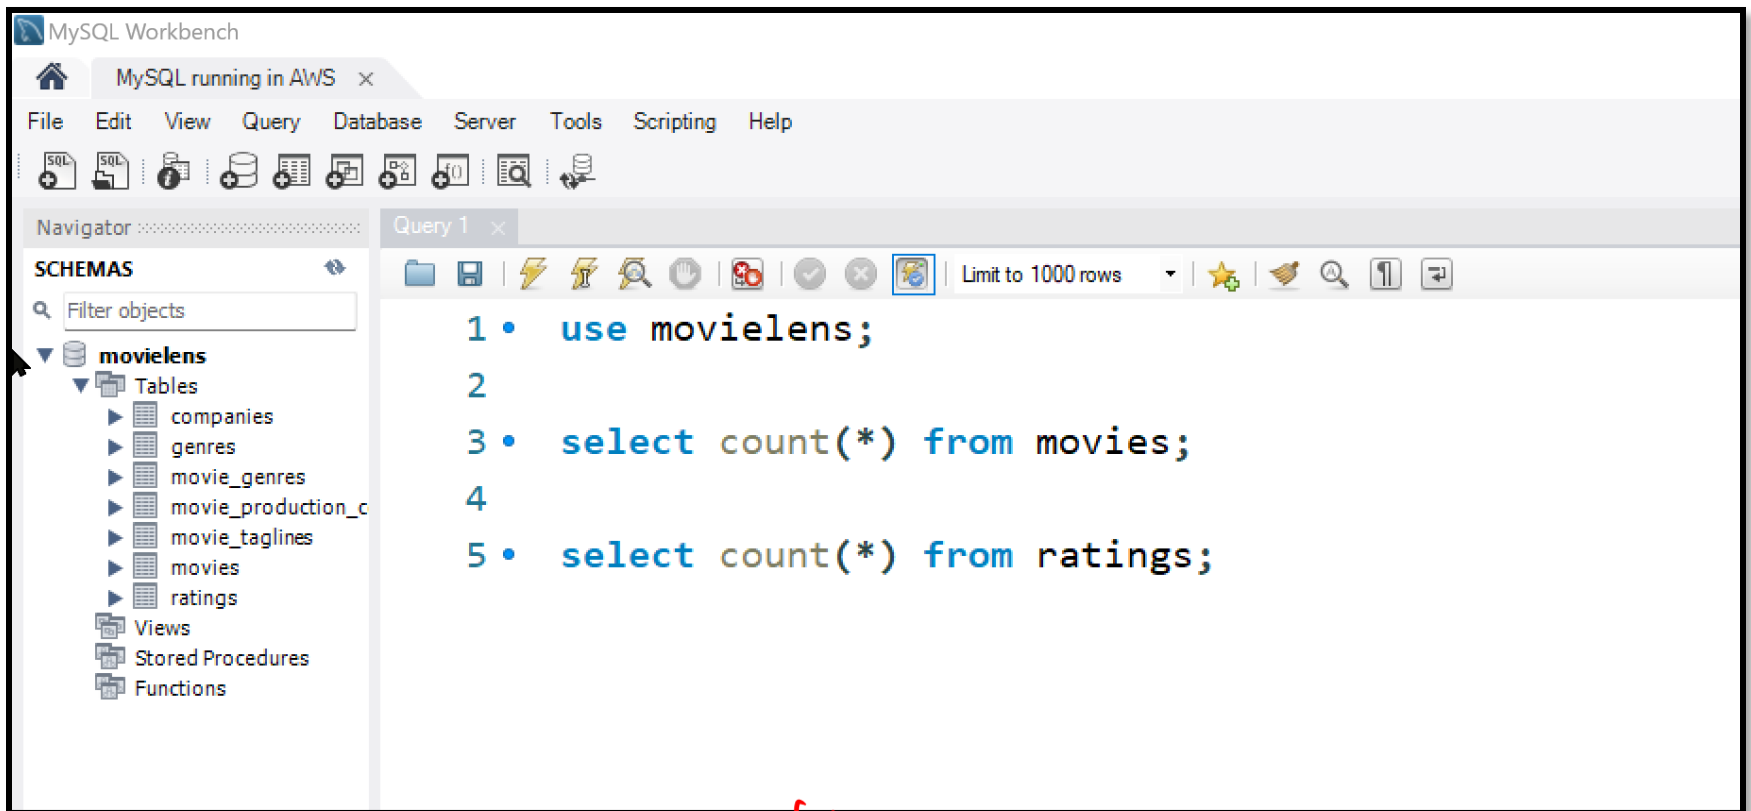
Username: movielens-read-only

Password: Store in Vault ... Clear

Default Schema:

Configure Server Management... Test Connection Cancel OK







**That's it, thank you!**



## Stations

Station_ID	Station_Name
40710	Chicago/Franklin
...	...

# CTA database

## Stops

Stop_ID	Station_ID	Stop_Name	Direction	ADA	Latitude	Longitude
30137	40710	Chicago (Kimball-Linden-bound)	N	1	41.89681	-87.635924
30138	40710	Chicago (Loop-bound)	S	1	41.89681	-87.635924
...	...	...	...	...	...	...

## Lines


Line_ID	Color
...	...
4	Brown
5	Purple
6	Purple-Express
...	...

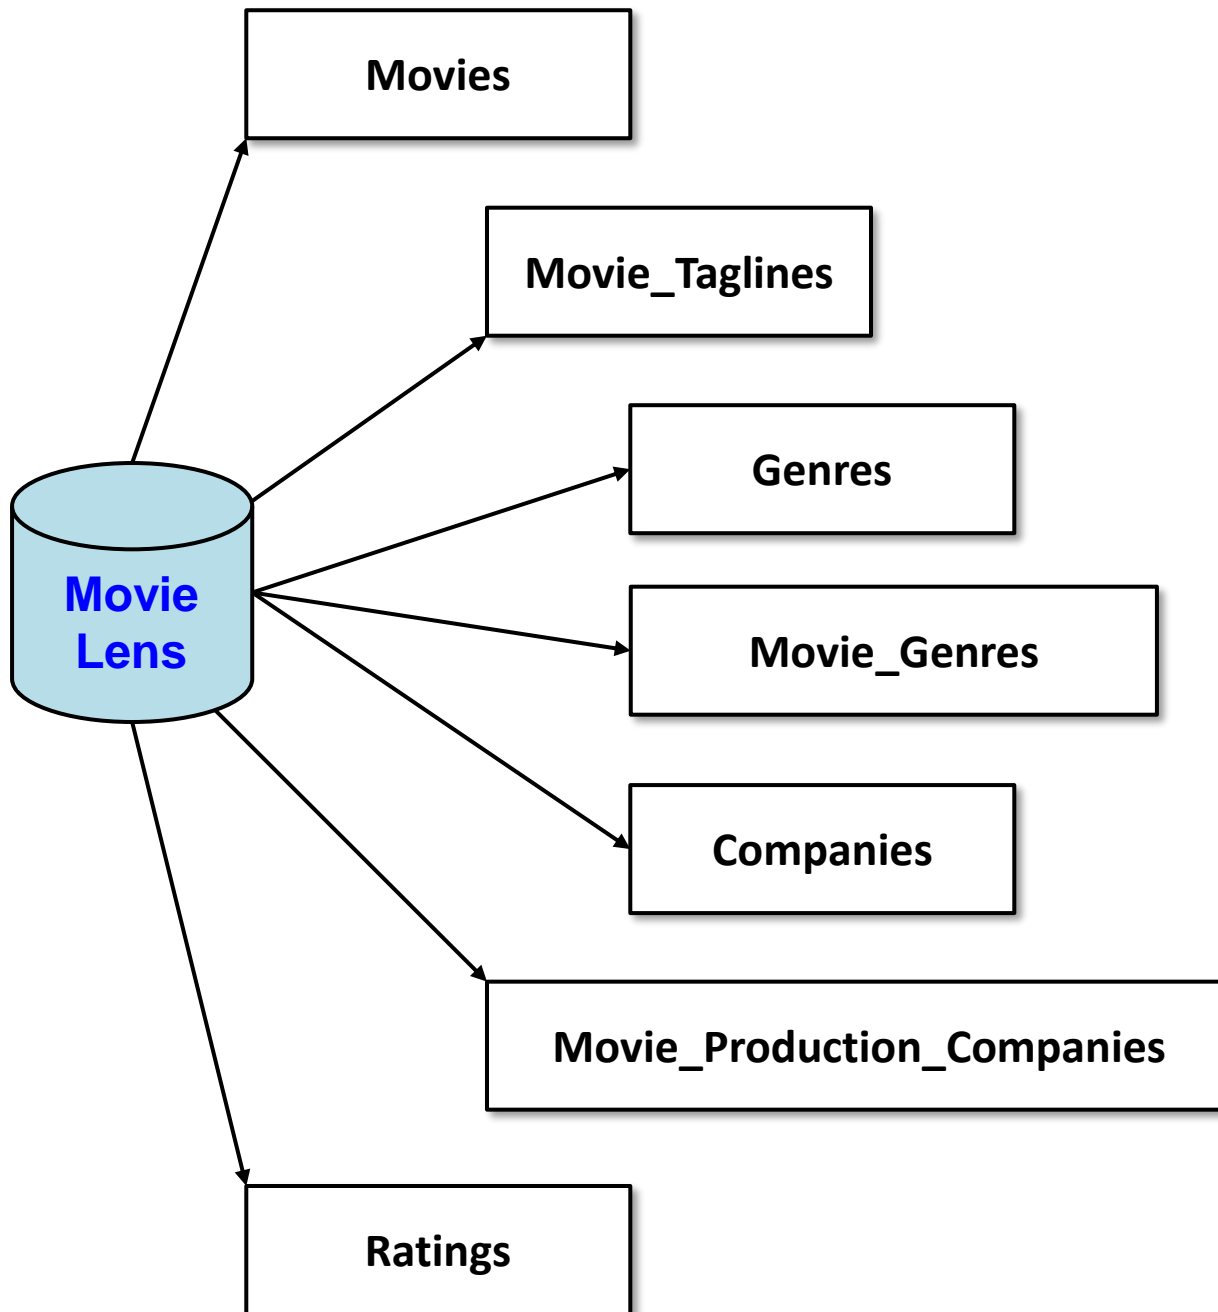
## LinesPerStop

Stop_ID	Line_ID
...	...
30137	4
30137	6
30138	4
30138	6
...	...

## Ridership


Station_ID	Ride_Date	Type_of_Day	Num_Riders
...	...	...	...
40710	2001-02-28 00:00:00.000	W	4206
...	...	...	...





Movie_ID	Title	Release_Date	Runtime	Original_L anguage	Budget	Revenue
603	The Matrix	1999-03-30 00:00:00.000	136	en	63000000	463517383
862	Toy Story	1995-10-30 00:00:00.000	81	en	30000000	373554033

Movies



Movie Taglines

Movie_ID	Tagline
603	Welcome to the Real World.
605	Everything that has a beginning has an end.

Ratings

Movie_ID	Rating
605	8
603	6
605	10
605	6

