

CS 310 : Scalable Software Architectures

Class session on Thursday, October 3rd



October 2024

| Sunday | Monday | Tuesday | Wednesday | Thursday | Friday | Saturday |
|--------|--------|---------|-----------|----------|--------|----------|
| | | 1 | 2 | 3 | 4 | 5 |
| 6 | 7 | 8 | 9 | 10 | 11 | 12 |
| 13 | 14 | 15 | 16 | 17 | 18 | 19 |
| 20 | 21 | 22 | 23 | 24 | 25 | 26 |
| 27 | 28 | 29 | 30 | 31 | | |

www.a-printable-calendar.com

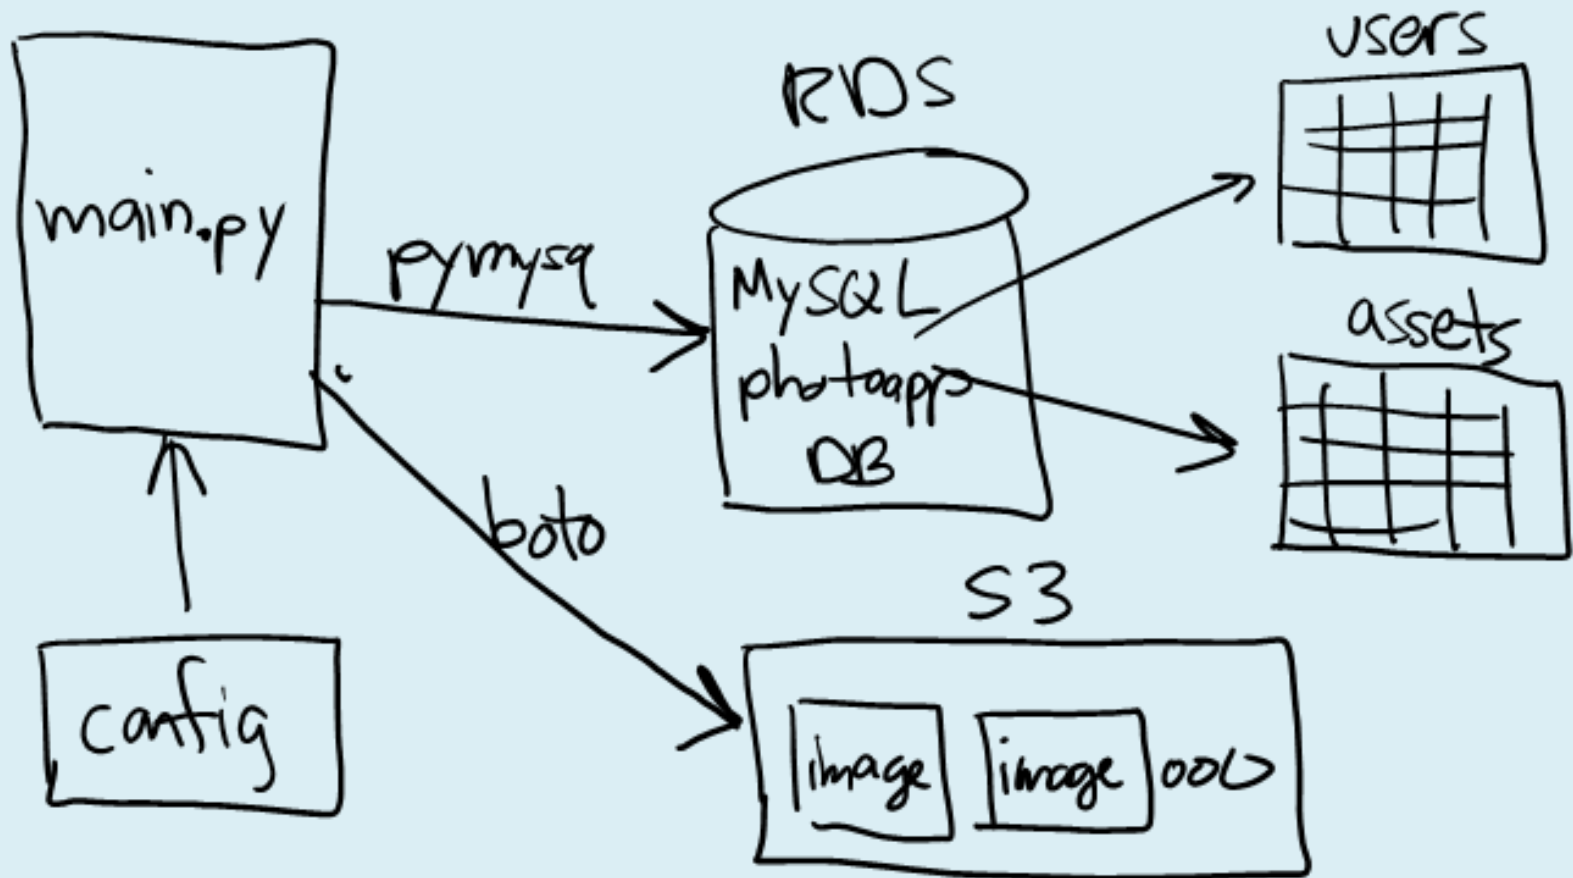
Notes:

- *Focus this week:*
 - *Relational databases*
- *Class sessions ***are*** being recorded this week*
 - *Will be available under Panopto on Canvas*
- *Project 01 due next Wednesday @ 11:59pm*
 - *Build a simple photo app using AWS*
 - *Parts 01 and 02 are both released*
 - *Can submit as late as Friday @ 11:59pm*
- *Office hours started Monday*
- *Optional SQL homework posted*



Northwestern
University

Project 01



Goals for today

- **Executing SQL from Python**

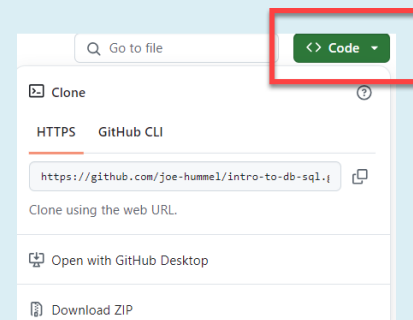
- *With `sqlite3`, a local file-based DBMS*
- *With `MySQL` running on a server in AWS*

- **Work with Docker**

- *Docker Desktop must be installed*
- *Download files from GitHub:*

- <https://github.com/joe-hummel/intro-to-db-sql>
- Clone repo or download ZIP


Same files as Tuesday,
no need to
clone/download again



Getting the necessary software

1. Download files you need for today

- <https://github.com/joe-hummel/intro-to-db-sql>
- Clone repo or download ZIP



2. Make sure Docker Desktop is running

3. Build Docker image and run container:

Linux/Mac/Windows WSL:

- 1) Open terminal, navigate to repo folder
- 2) `chmod 755 *.bash`
- 3) `./docker-build.bash`
- 4) `./docker-run.bash`

Windows:

- 1) Open Powershell, navigate to repo folder
- 2) `.\docker-build.bat`
- 3) `.\docker-run.bat`

```
hummel> ./docker-run.bash
docker> ls
Dockerfile      datatier.py      docker-build.bat  docker-run.bat    main.sql
cta.db          docker-build.bash  docker-run.bash   main.py           movielens.db
docker>
```

Common docker errors

1. "docker" command not found

- *Uninstall and reinstall Docker Desktop*

2. When you try to build, you are not authorized

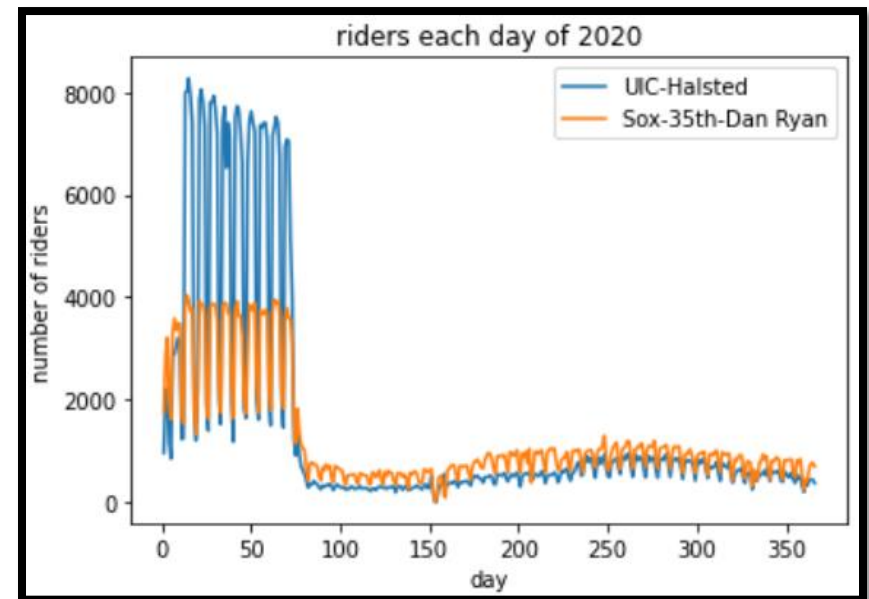
- *docker login -u docker-username*

3. When you try to run, you get errors like "bash: \$\r: command not found"


1. *If you see the **docker**> prompt, type **exit***
2. `((Get-Content .bashrc) -join "`n") + "`n" | Set-Content -NoNewLine .bashrc`

Executing SQL within other languages

- SQL is a powerful language, but lacking in other areas
 - *No UI support, no graphics, no web, no mobile*
- SQL is commonly embedded within other languages
 - *JavaScript*
 - *Python*
 - *Java*
 - *Swift*
 - *C#*
 - *etc.*



Executing SQL from Python



sqlite3 Python library

- **Open connection**

Cursor is a pointer to the data, and is used to scroll through the data

- **Create a cursor**

- **Execute SQL**

- **Fetch result**

- *fetchone()*

- Returns a tuple (...)

- *fetchall()*

- Returns a list of tuples

```
import sqlite3
```

```
dbConn = sqlite3.connect("filename.db")  
dbCursor = dbConn.cursor()
```

```
# query to retrieve 1 row from DB:
```

```
sql = """  
    Select ...  
    From ...  
    Where ... ;  
    """
```

```
dbCursor.execute(sql)  
row = dbCursor.fetchone()
```

```
print(row)
```

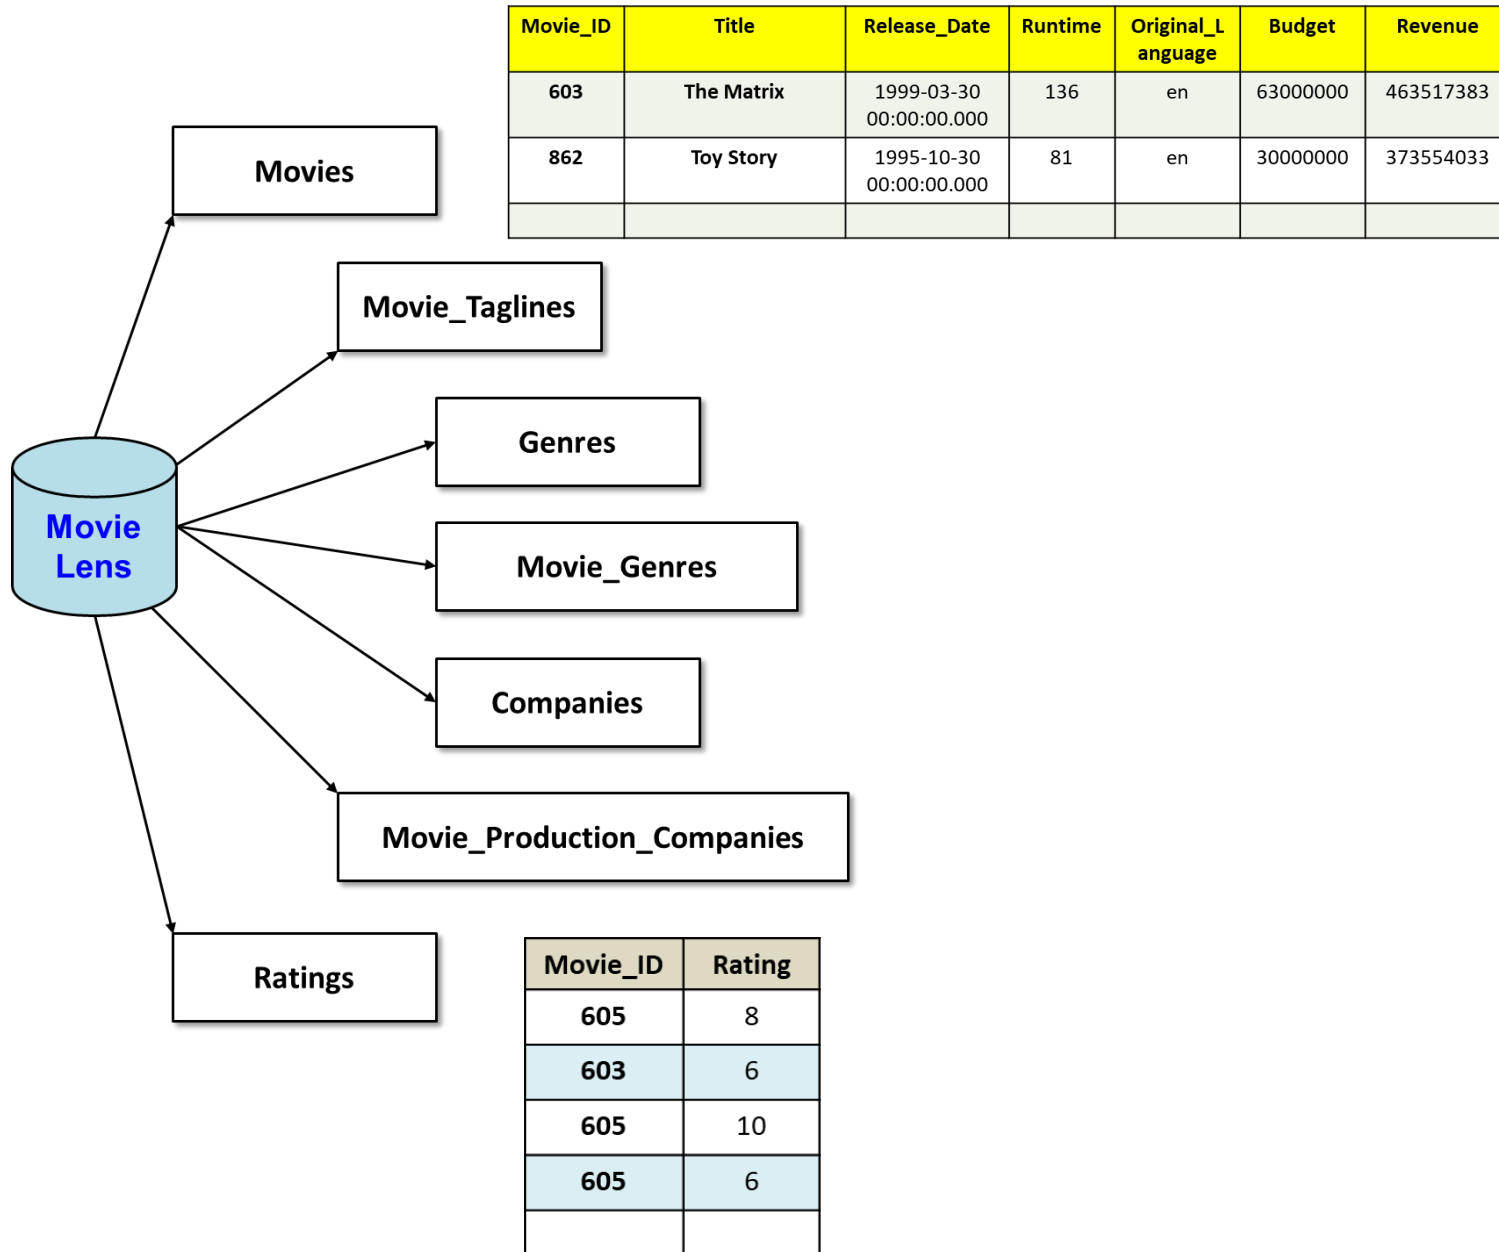
```
# query to retrieve N values from DB:
```

```
sql = """  
    Select ...  
    From ...  
    Where ... ;  
    """
```

```
dbCursor.execute(sql)  
rows = dbCursor.fetchall()
```

```
for row in rows:  
    print(row)
```


MovieLens database



Exercise #1

1. Build docker image (if you haven't already)
2. Run docker image
3. Create a file in the repo folder named "test.py"
4. Open "test.py" in your favorite text editor
5. Code in Python

- 1) *Open MovieLens.db*
- 2) *Retrieve all columns for 'Toy Story'*
- 3) *Print the row*

6. Docker> python3 test.py

```
docker> python3 test.py
(862, '1995-10-30 00:00:00.000', 81, 'en', 30000000, 373554033, 'Toy Story')
docker>
```

```
import sqlite3

dbConn = sqlite3.connect("movielens.db")
dbCursor = dbConn.cursor()

sql = """
    select * from movies
    where title = 'Toy Story';
    """

dbCursor.execute(sql)
row = dbCursor.fetchone()

print(row)
```

Your turn

- Output all movies that have > 200 reviews
 - *You need a join, group by, and having*
 - *Use fetchall()*

Movies

| Movie_ID | Title | Release_Date | Runtime | Original_L anguage | Budget | Revenue |
|----------|------------|----------------------------|---------|-----------------------|----------|-----------|
| 603 | The Matrix | 1999-03-30 00:00:00.000 | 136 | en | 63000000 | 463517383 |
| 862 | Toy Story | 1995-10-30 00:00:00.000 | 81 | en | 30000000 | |
| | | | | | | |

Ratings

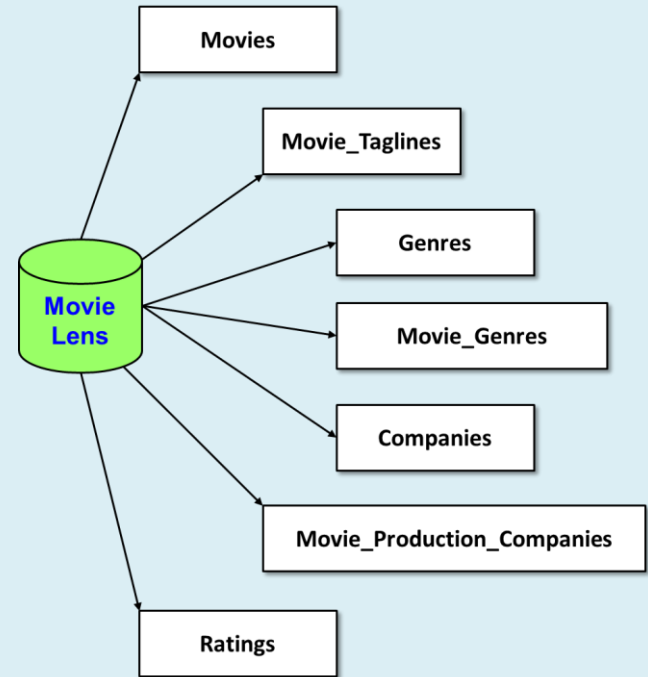
| Movie_ID | Rating |
|----------|--------|
| 605 | 8 |
| 603 | 6 |
| 605 | 10 |
| 605 | 6 |
| | |

```
docker> python3 test.py
License to Wed
Men in Black II
Monsoon Wedding
Once Were Warriors
Silent Hill
Sissi
Solaris
Terminator 3: Rise of the Machines
The 39 Steps
The Hours
The Million Dollar Hotel
The Passion of Joan of Arc
Three Colors: Red
docker>
```

Top-10 movies in Drama genre

Top-10 Drama movies

```
20 sql = ""
21     Select Title, Round(avg(Rating),2) as Rating
22     From Movies
23     Inner Join Ratings on Movies.Movie_ID = Ratings.Movie_ID
24     Inner Join Movie_Genres on Movies.Movie_ID = Movie_Genres.Movie_ID
25     Inner Join Genres on Genres.Genre_ID = Movie_Genres.Genre_ID
26     Where Genre_Name = 'Drama'
27     Group By Ratings.Movie_ID
28     Having Count(Rating) > 100
29     Order By Rating DESC, Title ASC
30     Limit 10;
31     ""
32
33 dbCursor.execute(sql)
34 rows = dbCursor.fetchall()
35
36 for row in rows:
37     print("Movie:", row[0], ", avg rating:", row[1]);
38
```



Movie: Sleepless in Seattle , avg rating: 8.98
Movie: The Million Dollar Hotel , avg rating: 8.97
Movie: Once Were Warriors , avg rating: 8.61
Movie: Confession of a Child of the Century , avg rating: 8.47
Movie: The Thomas Crown Affair , avg rating: 8.47
Movie: Scarface , avg rating: 8.45
Movie: Murder She Said , avg rating: 8.42
Movie: The Talented Mr. Ripley , avg rating: 8.4
Movie: Solaris , avg rating: 8.28
Movie: Notes on a Scandal , avg rating: 8.19

Execute the query

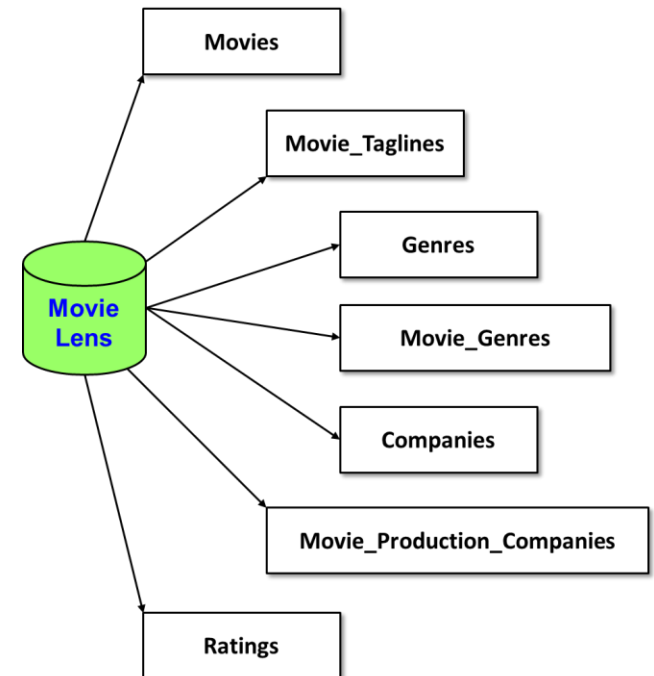
1. The code is already provided in "main.py"
2. Docker> python3 main.py

```
Movie: Sleepless in Seattle , avg rating: 8.98
Movie: The Million Dollar Hotel , avg rating: 8.97
Movie: Once Were Warriors , avg rating: 8.61
Movie: Confession of a Child of the Century , avg
rating: 8.47
Movie: The Thomas Crown Affair , avg rating: 8.47
Movie: Scarface , avg rating: 8.45
Movie: Murder She Said , avg rating: 8.42
Movie: The Talented Mr. Ripley , avg rating: 8.4
Movie: Solaris , avg rating: 8.28
Movie: Notes on a Scandal , avg rating: 8.19
```

Parameterized (“dynamic”) queries

- Most queries are dynamic, responding to user input

```
Enter genre: Comedy
('Sleepless in Seattle', 8.98)
('Men in Black II', 8.51)
('Shriek If You Know What I Did Last Friday the Thirteenth', 8.45)
('Murder She Said', 8.42)
('License to Wed', 8.36)
('Monsieur Hulot's Holiday', 8.28)
('The Prisoner of Zenda', 7.99)
('Fools Rush In', 7.97)
('My Name Is Bruce', 7.93)
('Blood: The Last Vampire', 7.91)
> □
```



Executing dynamic queries

```
import sqlite3
```

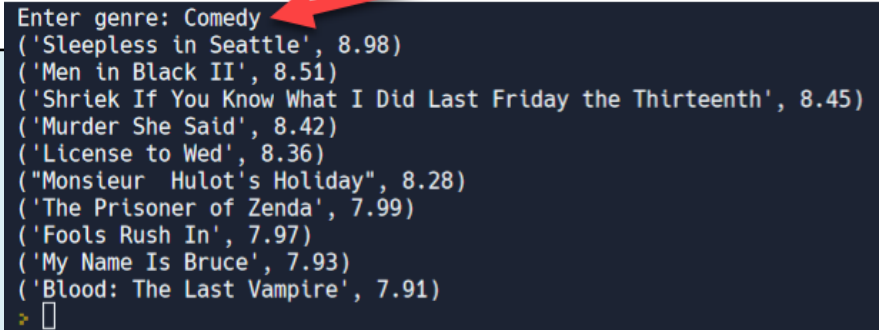
```
dbConn = sqlite3.connect("chicago-police-stops.db")  
dbCursor = dbConn.cursor()
```

```
genre = input("Enter genre: ")
```

```
sql = """  
    Select Title, Round(avg(Rating),2) as Rating  
    From    Movies  
    Inner Join Ratings on Movies.Movie_ID = Ratings.Movie_ID  
    Inner Join Movie_Genres on Movies.Movie_ID = Movie_Genres.Movie_ID  
    Inner Join Genres on Genres.Genre_ID = Movie_Genres.Genre_ID  
    Where Genre_Name = ?  
    Group By Movies.Movie_ID  
    Having Count(Rating) > 100  
    Order By Rating DESC, Title ASC  
    Limit 10;  
    """
```

```
dbCursor.execute(sql, [genre])  
rows = dbCursor.fetchall()
```

```
for row in rows:  
    print(row)
```



```
Enter genre: Comedy  
( 'Sleepless in Seattle', 8.98)  
( 'Men in Black II', 8.51)  
( 'Shriek If You Know What I Did Last Friday the Thirteenth', 8.45)  
( 'Murder She Said', 8.42)  
( 'License to Wed', 8.36)  
( 'Monsieur Hulot's Holiday', 8.28)  
( 'The Prisoner of Zenda', 7.99)  
( 'Fools Rush In', 7.97)  
( 'My Name Is Bruce', 7.93)  
( 'Blood: The Last Vampire', 7.91)  
> []
```

placeholder...

NOTE: with MySQL
use %s instead of ?

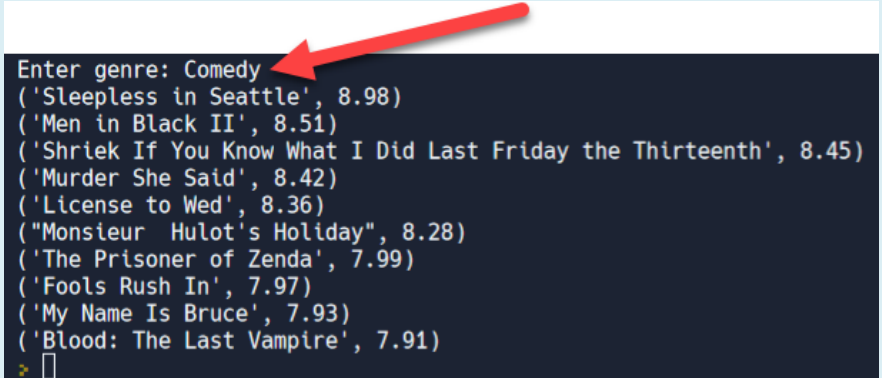
*Provide parameter value(s) in
call so DBMS builds query and
checks for injection attacks*

Try it

1. Modify "main.py" in your favorite editor

- 1) *Modify SQL query string, replacing 'Drama' with ?*
- 2) *Pass genre value in call to execute(sql, [genre])*
- 3) *Run and try different genres (Comedy, Horror, ...)*

2. Docker> python3 main.py



```
Enter genre: Comedy
('Sleepless in Seattle', 8.98)
('Men in Black II', 8.51)
('Shriek If You Know What I Did Last Friday the Thirteenth', 8.45)
('Murder She Said', 8.42)
('License to Wed', 8.36)
('Monsieur Hulot's Holiday', 8.28)
('The Prisoner of Zenda', 7.99)
('Fools Rush In', 7.97)
('My Name Is Bruce', 7.93)
('Blood: The Last Vampire', 7.91)
> 
```


Beware SQL injection attacks

```
genre = input("Enter genre: ")
```

```
sql = "..." + genre + "..."  
sql = f"... {genre} ..."
```

```
dbCursor.execute(sql)  
rows = dbCursor.fetchall()
```

```
for row in rows:  
    print(row)
```

NEVER build sql query strings yourself --- you open the door to SQL injection attacks

Attacks can delete your tables.


Attack Looks like this

ignore'; delete from ratings; —

ignore will ignore all the above
delete will remove all the data
— comments out everything after

These are very common attacks. So never compile SQL with f strings

<https://www.bleepingcomputer.com/news/security/researchers-find-sql-injection-to-bypass-airport-tsa-security-checks/>



How about executing SQL with MySQL?



`SELECT * FROM ...;`

Database
server



Database

pymysql

- Use pymysql module --- nearly identical API !
- <https://pymysql.readthedocs.io/en/latest/>

```
import pymysql

try:
    dbconn = pymysql.connect(
        host=ENDPOINT,
        port=PORTNUM,
        user=USERNAME,
        passwd=PASSWORD,
        database=DBNAME)

    sql = "..."

    dbCursor = dbconn.cursor()
    dbCursor.execute(sql)
    rows = dbCursor.fetchall()    # or fetchone()
    for row in rows:
        print(row)

except Exception as e:
    print("Database connection failed due to {}".format(e))
```

*For dynamic queries
use %s instead of ?*

Try it

1. Modify "Dockerfile" in your favorite editor

- 1) *On the last line add: `RUN pip3 install pymysql`*
- 2) *Exit your docker container: `exit`*
- 3) *Build docker image*
- 4) *Run docker image*

2. Download "main-mysql.py"

- 1) *<https://tinyurl.com/main-mysql>*
- 2) *Move "main-mysql.py" into your repo folder*

3. Docker> `python3 main-mysql.py`

Try it with these genres:

*Music
Family
Horror*

That's it, thank you!