

CS 310 : Scalable Software Architectures

Class session on Tuesday, October 22nd



October 2024

Sunday	Monday	Tuesday	Wednesday	Thursday	Friday	Saturday
		1	2	3	4	5
6	7	8	9	10	11	12
13	14	15	16	17	18	19
20	21	22	23	24	25	26
27	28	29	30	31		

www.a-printable-calendar.com

Notes:


- *Focus this week:*
 - *Scalability, performance, serverless*
- *Project 02 due Friday*
 - *Part 01: web service for photoapp, rewrite client*
 - *Part 02: deploy to AWS EC2*
 - *Due Friday with late submissions Sunday*
- *Exam next Tuesday*
 - *Practice problems and review in class on Thursday*



Northwestern
University

RESTful API

- Most web service APIs are said to be "RESTful"
- This is an architectural style with design constraints:




Constraints:

- Uniform Interface
- Stateless Operations
- Layered System
- Client-server based
- Cacheable
- Code on demand (optional)

that adheres to web server standards


Project 02 --- photoapp web service



```
//  
// PhotoApp web service  
//  
app.get('/stats', (req, res) => {...});  
app.put('/user', (req, res) => {...});  
app.get('/users', (req, res) => {...});  
app.get('/assets', (req, res) => {...});  
app.get('/bucket', (req, res) => {...});  
app.get('/image/:assetid', (req, res) => {...});  
app.post('/image/:userid', (req, res) => {...});
```

/user

- Design API to insert / update a user



– *Option #1: HTTP GET, sending data via parameters*

```
// data in URL /user?email=...&lastname=...&firstname=...  
app.get('/user', (req, res) => {...});
```

```
// data in URL /user/:email/:lastname/:firstname  
app.get('/user/:email/:last/:first', (req, res) => {...});
```

– *Option #2: HTTP PUT, sending data in body of request*

```
// put a new user or update an existing user  
app.put('/user', (req, res) => {...});
```

PUT makes more sense because we are inserting a user, not getting one. And we use PUT, not POST, because if repeated we want server to update not insert again.

/user -- server-side

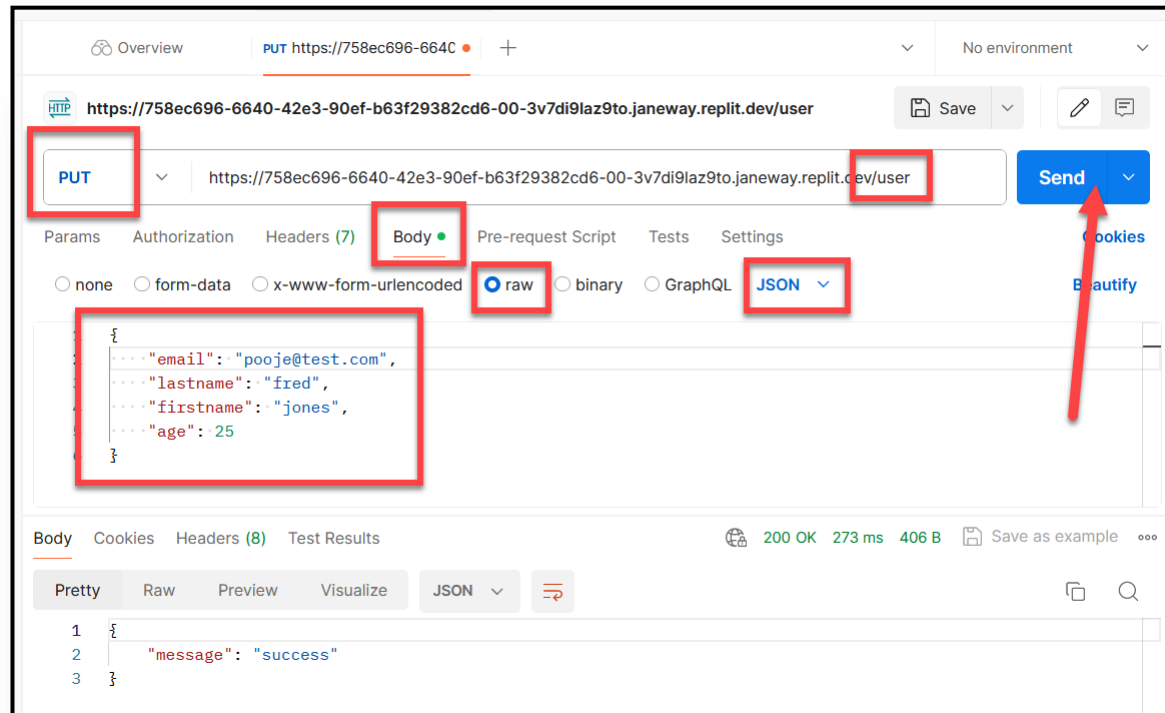


```
//  
// Insert/modify user in database:  
//  
app.use(express.json({strict: false})); // deserialize incoming JSON  
  
app.put('/user', (req, res) => {  
    console.log("**Call to /user...");  
  
    var data = req.body; // JSON data is deserialized => JS object  
  
    let email = data.email;  
  
    .  
    . // lookup email, if exists modify else insert  
    .  
  
    res.json(...); // send a response when done  
  
});
```



Testing with postman

- Browsers only GET
- Postman can GET / PUT / POST
 - <https://postman.com>



/user -- client-side


- The client **PUTs** user's data to the web service...

```
import requests

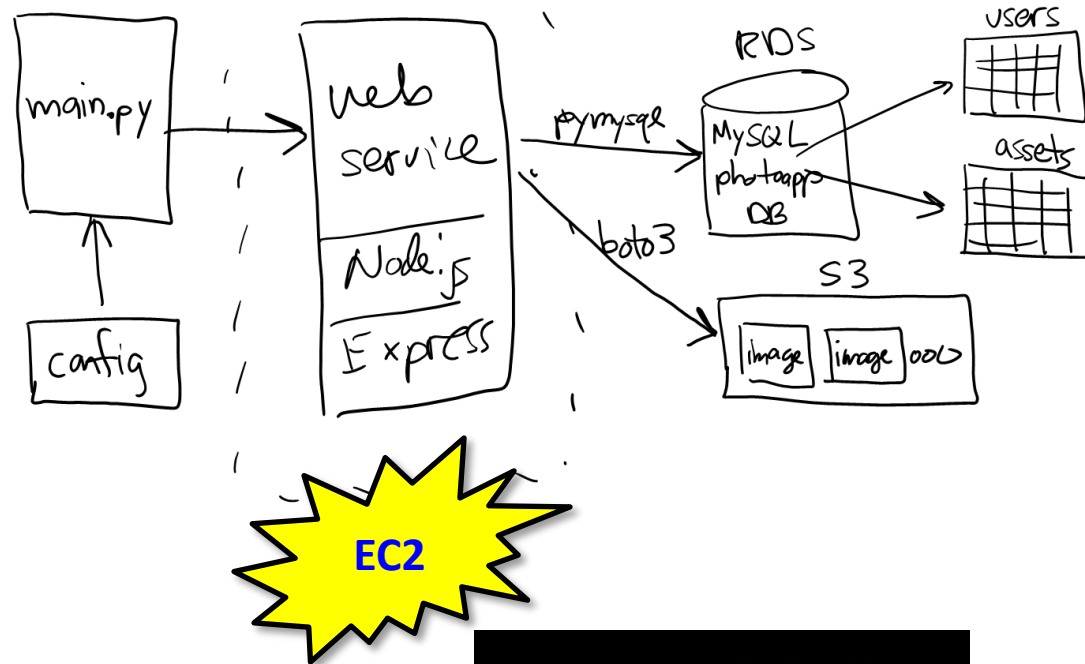
url = 'http://photoapp-web-service.com/user'

data = {
    "email":      "pooja@piazza.com",
    "lastname":   "sarkar",
    "firstname":  "pooja",
    "buckerfolder": "..."}

response = requests.put(url, json=data) # post data in JSON format
print("status code:", response.status_code)
```



Debugging in the cloud



EC2



Print debugging

- Print debugging is still very common --- view logs in AWS

```
app.put('/user', (req, res) => {  
  console.log("**Call to /user...");  
  var data = req.body;  
  console.log(data);  
  .  
  .  
  .  
})
```

Elastic Beanstalk > Environments > Nu-cs-msa-s3-photoapp-web-svc-env

Nu-cs-msa-s3-photoapp-web-svc-env Info



Actions ▼

Upload and deploy

Environment overview

Health

✓ Green

Domain

Nu-cs-msa-s3-photoapp-web-svc-env.eba-njwcmj6i.us-east-2.elasticbeanstalk.com

Environment ID

e-cdg9y4f7ma

Application name

nu-cs-msa-s3-photoapp-web-service

Platform

Change version

Platform

Node.js 18 running on 64bit Amazon Linux 2/5.9.7

Running version

photoapp-web-service-v1

Platform state

✓ Supported

Events

Health

Logs

Monitoring

Alarms

Managed updates

Tags

Logs Info



Request logs ▼




/var/log/web.stdout.log

```
Oct 22 05:25:50 ip-172-30-0-130 web: /stats: calling S3...
Oct 22 05:25:50 ip-172-30-0-130 web: /stats: calling RDS to get # of users...
Oct 22 05:25:50 ip-172-30-0-130 web: /stats: calling RDS to get # of assets...
Oct 22 05:25:50 ip-172-30-0-130 web: /stats done, sending response...
Oct 22 05:25:50 ip-172-30-0-130 web: **Call to get /stats...
Oct 22 05:25:50 ip-172-30-0-130 web: /stats: calling S3...
Oct 22 05:25:50 ip-172-30-0-130 web: /stats: calling RDS to get # of users...
Oct 22 05:25:50 ip-172-30-0-130 web: /stats: calling RDS to get # of assets...
Oct 22 05:25:50 ip-172-30-0-130 web: /stats done, sending response...
Oct 22 05:25:50 ip-172-30-0-130 web: **Call to get /stats...
Oct 22 05:25:50 ip-172-30-0-130 web: /stats: calling S3...
Oct 22 05:25:50 ip-172-30-0-130 web: /stats: calling RDS to get # of users...
Oct 22 05:25:50 ip-172-30-0-130 web: /stats: calling RDS to get # of assets...
Oct 22 05:25:50 ip-172-30-0-130 web: /stats done, sending response...
Oct 22 05:25:50 ip-172-30-0-130 web: **Call to get /stats...
Oct 22 05:25:50 ip-172-30-0-130 web: /stats done, sending response...
Oct 22 05:25:50 ip-172-30-0-130 web: **Call to get /stats...
Oct 22 05:25:50 ip-172-30-0-130 web: /stats: calling S3...
Oct 22 05:25:50 ip-172-30-0-130 web: /stats: calling RDS to get # of users...
Oct 22 05:25:50 ip-172-30-0-130 web: /stats: calling RDS to get # of assets...
Oct 22 05:25:50 ip-172-30-0-130 web: /stats done, sending response...
Oct 22 05:25:50 ip-172-30-0-130 web: **Call to get /stats...
Oct 22 05:25:50 ip-172-30-0-130 web: /stats: calling S3...
Oct 22 05:25:50 ip-172-30-0-130 web: /stats: calling RDS to get # of users...
Oct 22 05:25:50 ip-172-30-0-130 web: /stats: calling RDS to get # of assets...
Oct 22 05:25:50 ip-172-30-0-130 web: /stats done, sending response...
```

Performance of cloud-based apps

- **Latency**: how long does it take? (speed = response time)
- **Bandwidth**: how much data can be transmitted? (volume)
- **Throughput**: how much processing per time unit? (# of users)



ab == apache benchmark

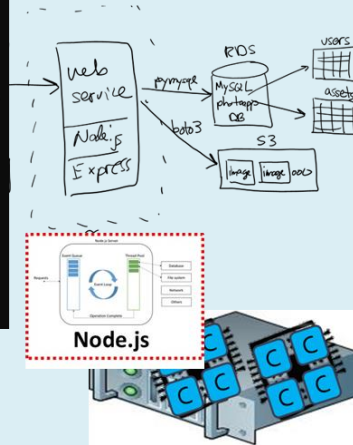
```
humme1> ab -c 1 -n 10 http://photoapp-nu-web-service-env-2.eba-htg5fauw.us-east-2.elasticbeanstalk.com/
This is ApacheBench, Version 2.3 <$Revision: 1643412 $>
Copyright 1996 Adam Twiss, Zeus Technology Ltd, http://www.zeustech.net/
Licensed to The Apache Software Foundation, http://www.apache.org/

Benchmarking photoapp-nu-web-service-env-2.eba-htg5fauw.us-east-2.elasticbeanstalk.com (be patient).....done

Server Software:      nginx
Server Hostname:      photoapp-nu-web-service-env-2.eba-htg5fauw.us-east-2.elasticbeanstalk.com
Server Port:          80


Document Path:        /
Document Length:      74 bytes

Concurrency Level:    1
Time taken for tests:  0.617 seconds
Complete requests:    10
```




- Reference: <https://httpd.apache.org/docs/2.4/programs/ab.html>
 - Linux: `sudo apt-get install apache2-tools`
 - Mac: already installed (terminal window)
 - Windows: download apache (<https://www.apachelounge.com/download/>), extract `/bin/ab.exe`
- Usage: **ab -k -c 10 -n 1000 URL**
 - k => keep-alive the TCP connection (cold-start vs. warm-start)
 - c => concurrency level (use this to simulate concurrent users / load on the server)
 - n => # of requests (use this to get a more accurate "average" response time)

Question




- Baseline latency is _____ ms
- Speed of RDS _____ ms
- Speed of S3 _____ ms






```
//  
// PhotoApp web service  
//  
app.get('/stats', (req, res) => {...});  
app.get('/users', (req, res) => {...});  
app.get('/assets', (req, res) => {...});  
app.get('/bucket', (req, res) => {...});  
app.get('/image/:assetid', (req, res) => {...});
```

Question




- Throughput...
- Node.js is single-threaded and runs on one core. How many users can it support before latency suffers?




```
//  
// PhotoApp web service  
//  
app.get('/stats', (req, res) => {...});  
app.get('/users', (req, res) => {...});  
app.get('/assets', (req, res) => {...});  
app.get('/bucket', (req, res) => {...});  
app.get('/image/:assetid', (req, res) => {...});
```

Question



- How efficient is the network (bandwidth)?
- Asset 1004 is 43K, asset 1014 is 1,964K --- 46x bigger
 - Time to download image # 1004 is _____ ms
 - How much longer does it take to download image # 1014? (_____ ms)



```
//  
// PhotoApp web service  
//  
app.get('/stats', (req, res) => {...});  
app.get('/users', (req, res) => {...});  
app.get('/assets', (req, res) => {...});  
app.get('/bucket', (req, res) => {...});  
app.get('/image/:assetid', (req, res) => {...});
```

That's it, thank you!