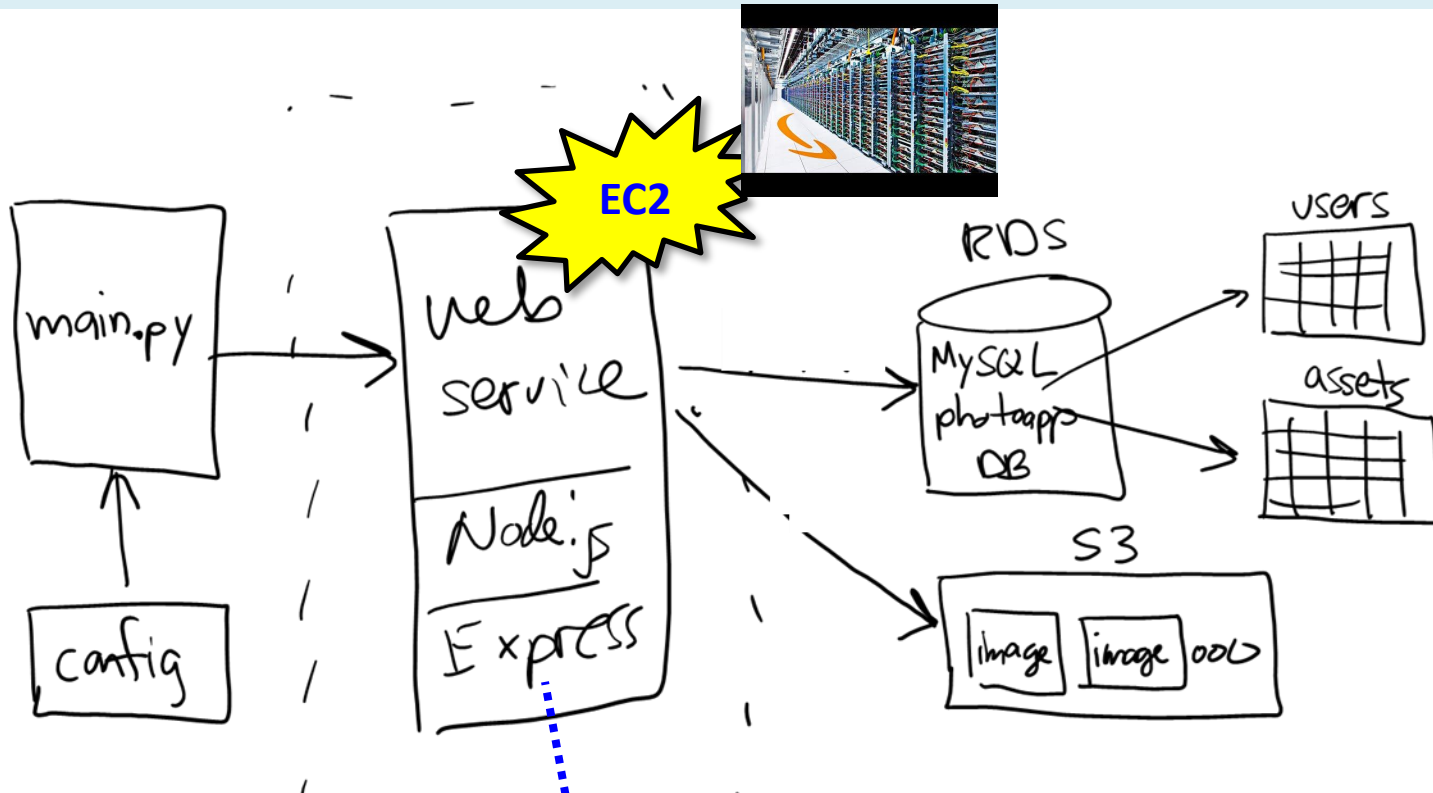


Performance

- **Performance in the cloud**
- **Latency, bandwidth, throughput**
- **Measuring performance with apache benchmark**







Project 02 --- photoapp web service in AWS



```
//  
// PhotoApp web service  
//  
app.get('/stats', (req, res) => {...});  
app.put('/user', (req, res) => {...});  
app.get('/users', (req, res) => {...});  
app.get('/assets', (req, res) => {...});  
app.get('/bucket', (req, res) => {...});  
app.get('/image/:assetid', (req, res) => {...});  
app.post('/image/:userid', (req, res) => {...});
```


Demo

<http://nu-cs-msa-s3-photoapp-web-servic-env.eba-njwcmj6i.us-east-2.elasticbeanstalk.com/stats>



Network latency vs. bandwidth

- **Latency**: how long does it take? (speed = response time)
- **Bandwidth**: how much data can be transmitted? (volume)
- **Throughput**: how much processing per time unit? (# of users)



Benchmarking

- Let's benchmark project 02 web service
- Baseline:

– *cost of a round-trip with no computation*

=> time to retrieve "home page"...

<http://nu-cs-msa-s3-photoapp-web-service-env.eba-njwcmj6i.us-east-2.elasticbeanstalk.com/>




```
hummel> ab -c 1 -n 10 http://photoapp-nu-web-service-env-2.eba-htg5fauw.us-east-2.elasticbeanstalk.com/
This is ApacheBench, Version 2.3 <Revision: 1843412>
Copyright 1996 Adam Twiss, Zeus Technology Ltd, http://www.zeustech.net/
Licensed to The Apache Software Foundation, http://www.apache.org/

Benchmarking photoapp-nu-web-service-env-2.eba-htg5fauw.us-east-2.elasticbeanstalk.com (be patient).....done

Server Software:      nginx
Server Hostname:      photoapp-nu-web-service-env-2.eba-htg5fauw.us-east-2.elasticbeanstalk.com
Server Port:          80

Document Path:        /
Document Length:      74 bytes

Concurrency Level:    1
Time taken for tests:  0.617 seconds
Complete requests:    10
```



ab == apache benchmark

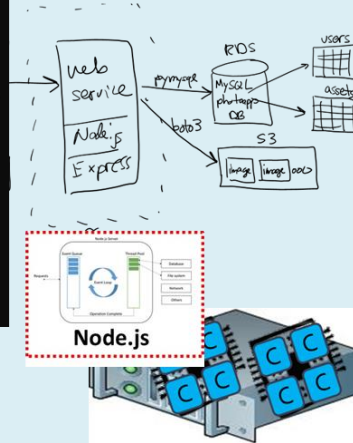
```
humme1> ab -c 1 -n 10 http://photoapp-nu-web-service-env-2.eba-htg5fauw.us-east-2.elasticbeanstalk.com/
This is ApacheBench, Version 2.3 <$Revision: 1643412 $>
Copyright 1996 Adam Twiss, Zeus Technology Ltd, http://www.zeustech.net/
Licensed to The Apache Software Foundation, http://www.apache.org/

Benchmarking photoapp-nu-web-service-env-2.eba-htg5fauw.us-east-2.elasticbeanstalk.com (be patient).....done

Server Software:      nginx
Server Hostname:      photoapp-nu-web-service-env-2.eba-htg5fauw.us-east-2.elasticbeanstalk.com
Server Port:          80


Document Path:        /
Document Length:      74 bytes

Concurrency Level:    1
Time taken for tests:  0.617 seconds
Complete requests:    10
```




- Reference: <https://httpd.apache.org/docs/2.4/programs/ab.html>
 - Linux: *sudo apt-get install apache2-tools*
 - Mac: *already installed (terminal window)*
 - Windows: *download apache (<https://www.apachelounge.com/download/>), extract /bin/ab.exe*
- Usage: **ab -k -c 10 -n 1000 URL**
 - k => *keep-alive the TCP connection (cold-start vs. warm-start)*
 - c => *concurrency level (use this to simulate concurrent users / load on the server)*
 - n => *# of requests (use this to get a more accurate "average" response time)*

Question #1



- Baseline latency is _____ ms
- Best guess --- which API function(s) will have fastest response time / lowest latency?




```
//  
// PhotoApp web service  
//  
app.get('/stats', (req, res) => {...});  
app.get('/users', (req, res) => {...});  
app.get('/assets', (req, res) => {...});  
app.get('/bucket', (req, res) => {...});  
app.get('/image/:assetid', (req, res) => {...});
```



Observations...

1. **/users** and **/assets** are fastest
2. **Why? RDS is faster than S3**
 - *DB vs. Web service*
3. **Notice that /users latency is close to baseline**
 - *The trip to AWS is the largest part of the cost...*
 - *Remember to always minimize the # of trips...*


Question #2




- Does async programming make a difference on the server?



Question #3



- Let's look at bandwidth vs. latency...
- Asset 1004 is 43K, asset 1014 is 1,964K --- 46x bigger
 - Time to download image # 1004 is _____ ms
 - How much longer will it take to download image # 1014? (_____ ms)




```
//  
// PhotoApp web service  
//  
app.get('/stats', (req, res) => {...});  
app.get('/users', (req, res) => {...});  
app.get('/assets', (req, res) => {...});  
app.get('/bucket', (req, res) => {...});  
app.get('/image/:assetid', (req, res) => {...});
```




Observations...

- **You can generally ignore payload size, and focus on minimizing the # of trips**
 - *Downloading an image doesn't take much longer than other functions*
 - *Even though #1007 is 40x larger, only takes 3x time*
- ==> not faster to break apart and send separately*

Question #4



- Throughput...
- Node.js is single-threaded and runs on one core. How many users can it support before latency suffers?
 - use *ab -c* option




```
//  
// PhotoApp web service  
//  
app.get('/stats', (req, res) => {...});  
app.get('/users', (req, res) => {...});  
app.get('/assets', (req, res) => {...});  
app.get('/bucket', (req, res) => {...});  
app.get('/image/:assetid', (req, res) => {...});
```

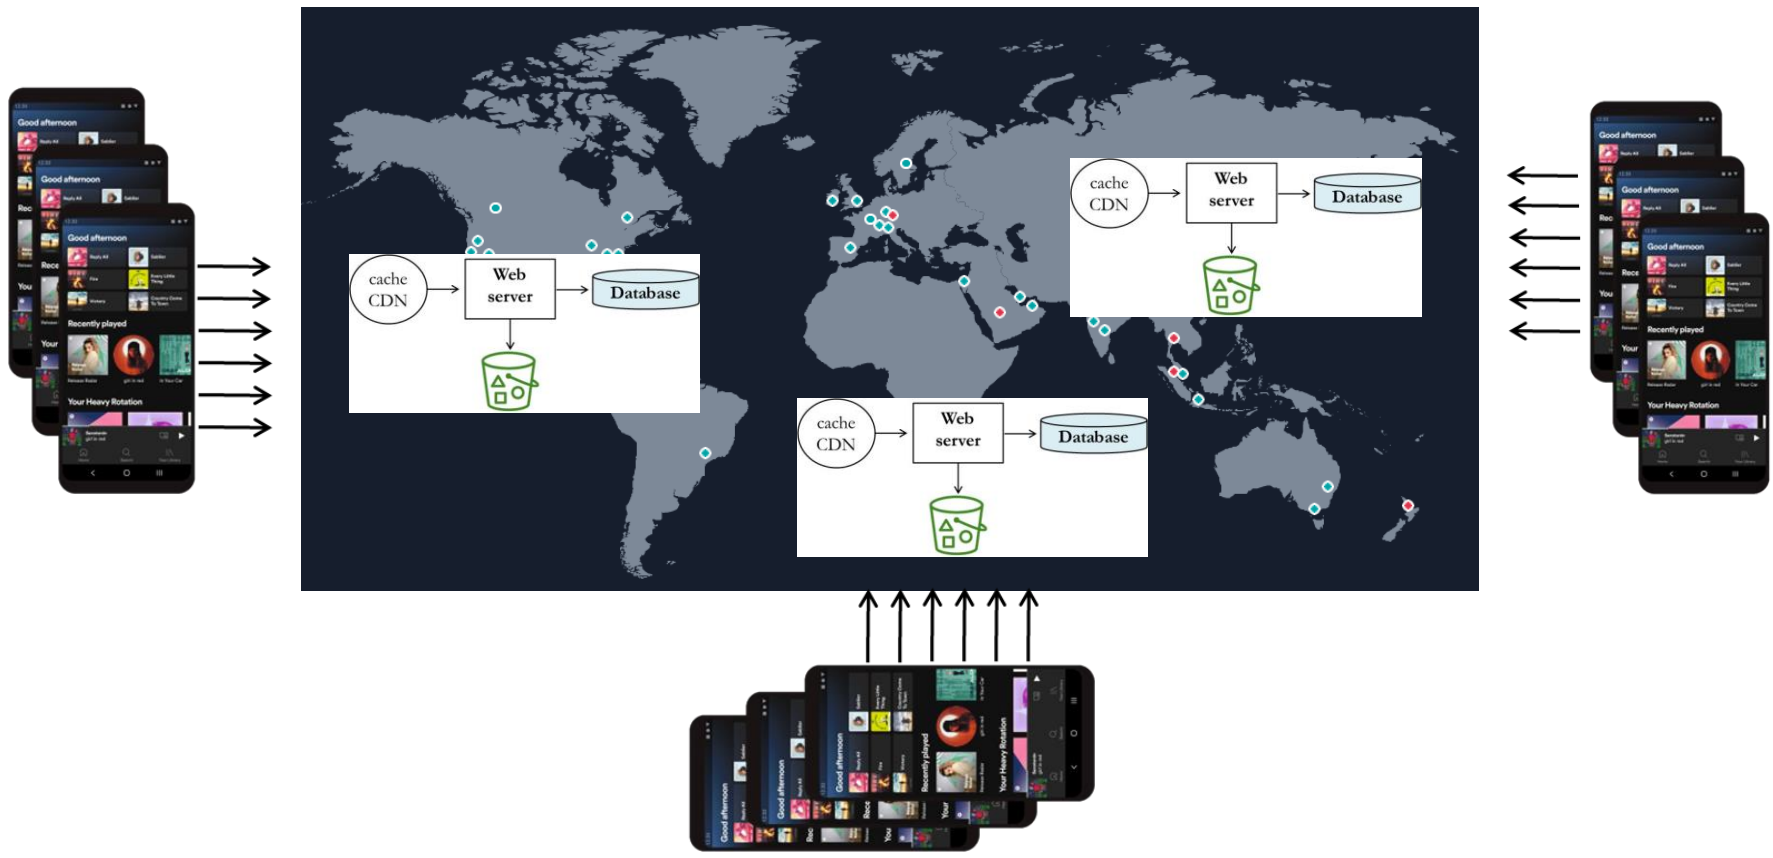
Observations...

- 1. Definitely more than one user...**
 - *Due to asynchronous programming model*
- 2. Typically 5-10 users depending on the function**

How to improve latency?




- Faster internet connection or shorten the distance...
 - Cache data in CDN (AWS CloudFront)
 - Replicate in different regions (USA, Europe, Asia, Africa, ...)




How to improve bandwidth?

- Bandwidth is a function of the network connection
- You generally don't have much control
 - *Mobile, Wifi, hard-wired, ...*
 - *All different speeds with different reliability*



How to improve throughput (# of users)?


Vertical scaling makes your machine(s) bigger and stronger.
Think more cores, RAM.



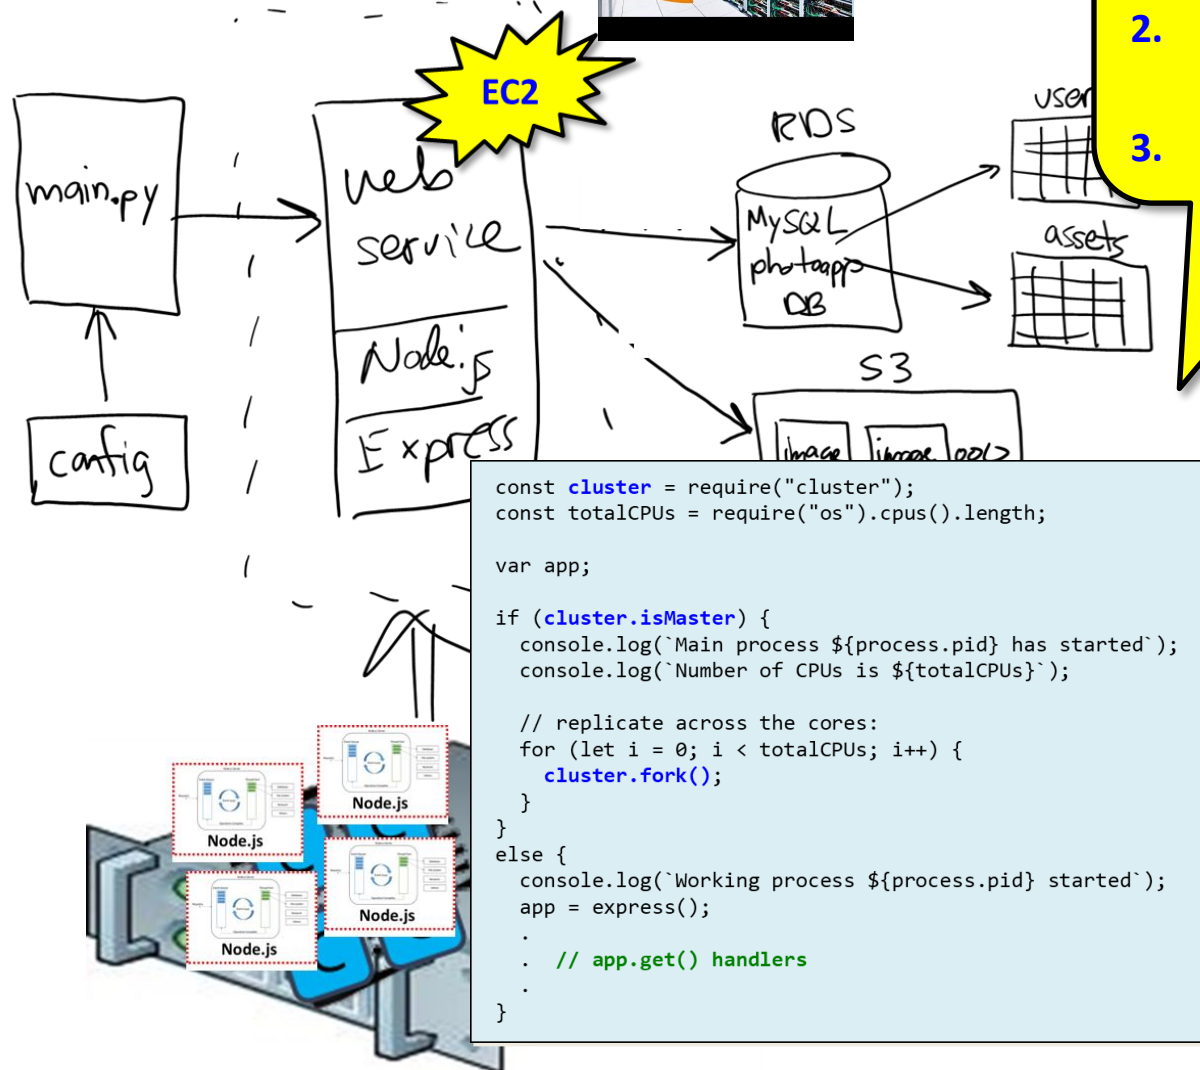
Horizontal scaling adds more machines.
Think of them standing side-by-side.

- **Vertical scaling** is easy, rent more cores / RAM
 - *Instant scalability, but cores may sit idle most of the time (wasting \$)*
- **Horizontal scaling** supports any # of users
 - *May take a few minutes for machines to startup (some users wait)*

- Throughput improved by scaling



Scaling up project 02 on EB / EC2?



1. Rent a VM with more cores & RAM
2. Rewrite app.js to fork / replicate
3. Deploy new version

That's it, thank you!