

CS 310 : Scalable Software Architectures

Class session on Tuesday, November 12th



NOVEMBER 2024

Sunday	Monday	Tuesday	Wednesday	Thursday	Friday	Saturday
					1	2
3	4	5	6	7	8	9
10	11	12	13	14	15	16
17	18	19	20	21	22	23
24	25	26	27	28	29	30

www.a-printable-calendar.com

Notes:

- *Focus this week:*
 - *Software execution and packaging*
 - *Docker*
 - *Async architectures, Distributed systems*
 - *CAP theorem*
- *Midterms have been graded; details in email*
- *Project 03 due Friday (Sunday late)*
 - *Serverless, event-driven PDF analysis*
 - *Gradescope will open today*
- *No class session Thursday --- office hours instead*
- *No class Tuesday 11/26 --- thanksgiving week*




Northwestern
University

Getting the necessary software

1. Make sure Docker Desktop is running

2. Download files you need for today

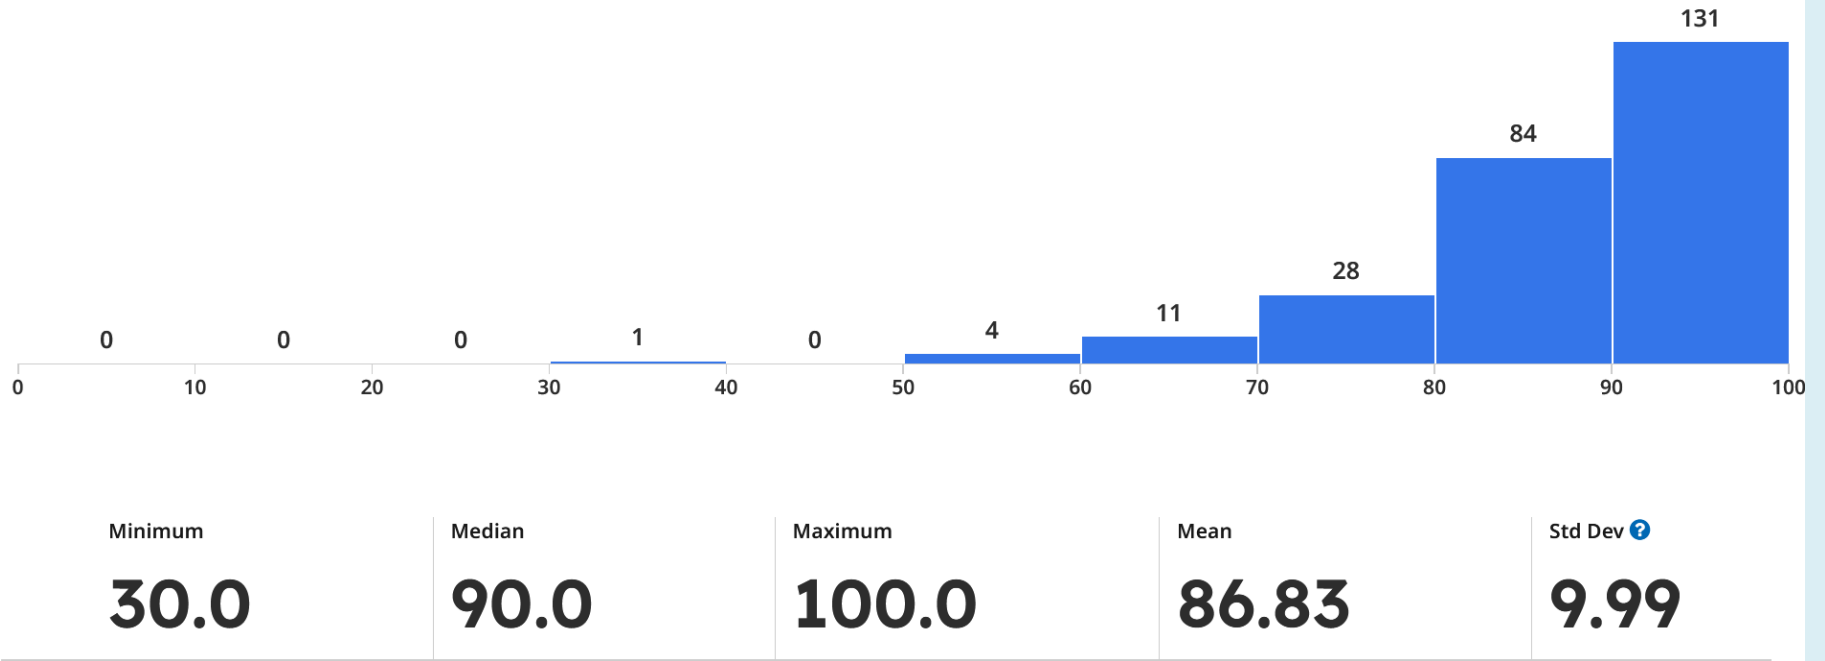
- <https://github.com/joe-hummel/docker-lambda-movies>



• Install desktop version of Postman app

- visit <https://www.postman.com/>
- download desktop app (click link in lower-left corner)

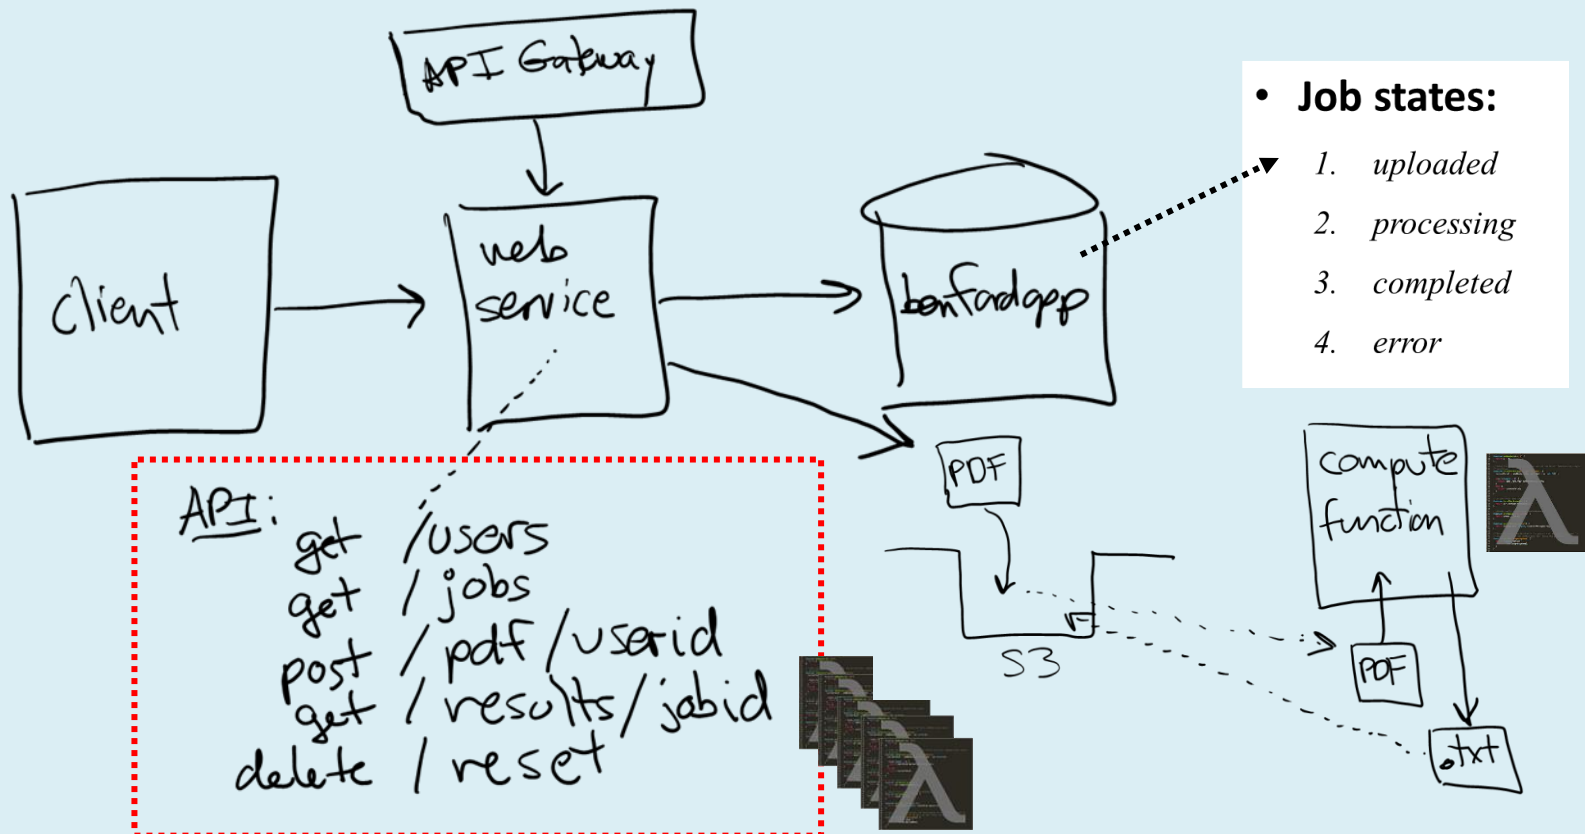
Midterm



Project 03

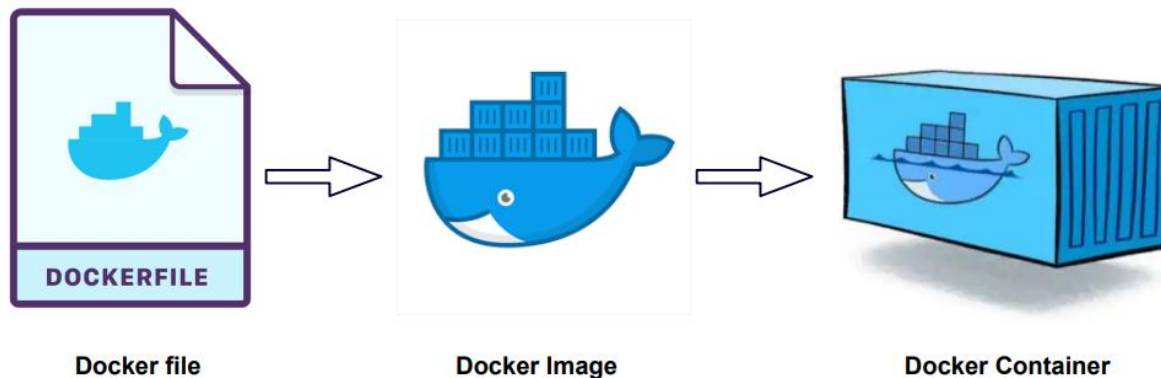
- Project 03 uses a different architectural approach:

1. *Event-driven*
2. *Serverless*
3. *Asynchronous API*



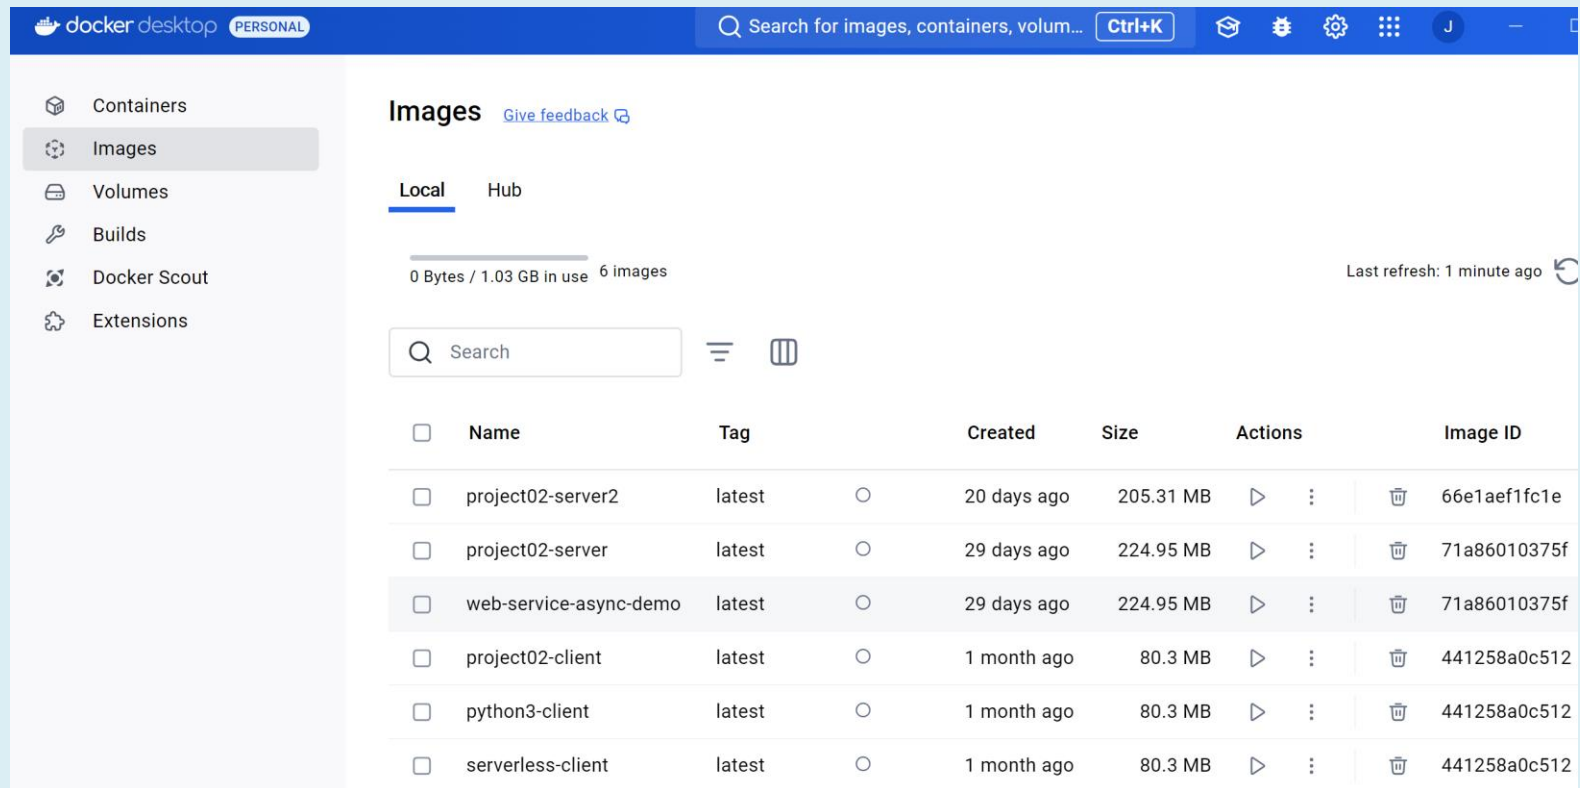
Overview

- A **Dockerfile** is a text document that contains the commands needed to build a **docker image**
- Think of a **docker image** as a snapshot of the software – a large file that can be stored and shared
 - *Docker Hub* is the app store for docker images
- A **docker container** runs a docker image



Images vs. Containers

- **Docker desktop** will show your images and containers...



The screenshot shows the Docker Desktop interface. The left sidebar contains navigation links: Containers, Images (selected), Volumes, Builds, Docker Scout, and Extensions. The main panel is titled 'Images' and shows the 'Local' tab. It displays a summary: '0 Bytes / 1.03 GB in use' and '6 images'. Below this is a search bar and a table of local images.

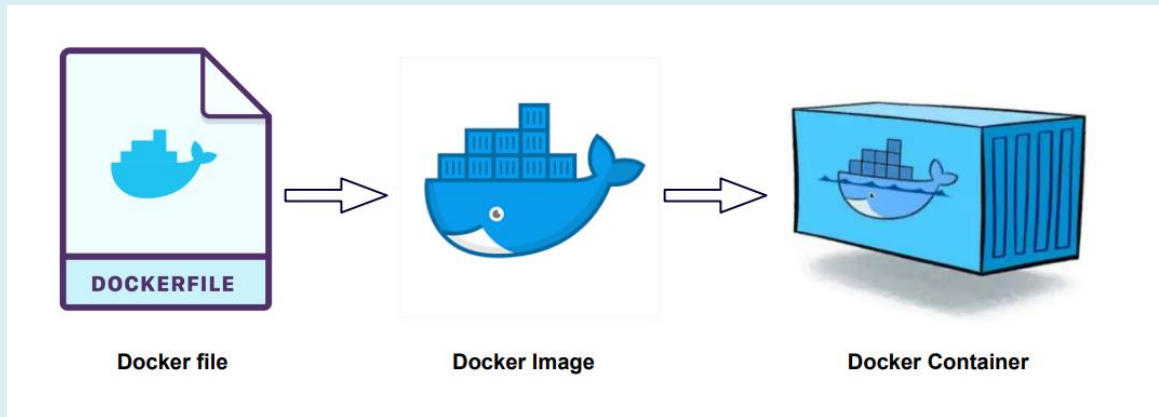
<input type="checkbox"/>	Name	Tag		Created	Size	Actions	Image ID
<input type="checkbox"/>	project02-server2	latest	○	20 days ago	205.31 MB	▶ ⋮	66e1aef1fc1e
<input type="checkbox"/>	project02-server	latest	○	29 days ago	224.95 MB	▶ ⋮	71a86010375f
<input type="checkbox"/>	web-service-async-demo	latest	○	29 days ago	224.95 MB	▶ ⋮	71a86010375f
<input type="checkbox"/>	project02-client	latest	○	1 month ago	80.3 MB	▶ ⋮	441258a0c512
<input type="checkbox"/>	python3-client	latest	○	1 month ago	80.3 MB	▶ ⋮	441258a0c512
<input type="checkbox"/>	serverless-client	latest	○	1 month ago	80.3 MB	▶ ⋮	441258a0c512

Docker Commands

Command	Usage	Example
docker run	Start a new container	<code>docker run image_name</code>
docker build	Build a Docker image from a Dockerfile	<code>docker build -t image_name .</code>
docker pull	Pull an image or repository from Docker registry	<code>docker pull ubuntu</code>
docker push	Push a local Docker image to a registry	<code>docker push my_username/image_name</code>
docker images	List all the Docker images stored locally	<code>docker images</code>
docker ps	Shows all running containers. Flag -a shows all containers (running & stopped).	<code>docker ps -a</code>
docker stop	Stops one or more running containers	<code>docker stop container_name</code>
docker rm	Remove a specific stopped container	<code>docker rm container_id</code>
docker rmi	Remove one or more Docker images	<code>docker rmi image_name</code>

Example #1

- Let's run MySQL on your local machine using Docker...



Step #1

1. Create the Dockerfile

```
Dockerfile
1 FROM mysql:latest
2 #
3 # set root password for local execution:
4 #
5 ENV MYSQL_ROOT_PASSWORD=abc123
6 #
7 # NOTE: changing to port 3307 since I already have MySQL installed
8 # and running locally on its own. So this docker image is a second
9 # version.
10 #
11 ENV MYSQL_TCP_PORT=3307
12 #
13 # expose the port needed to connect to MySQL server:
14 #
15 EXPOSE 3307
16 #
17 # start server when container runs:
18 #
19 CMD ["mysqld"]
20
```

Step #2

2. Build the docker image “mysql”

- t tag the image so easier to refer to*
- . refer to Dockerfile in current directory*

```
docker build -t mysql .
```

```
docker images
```

Step #3

3. Run the docker image => docker container

-d detach from terminal window to run in background

-p map local port to container port

--name container so easier to refer to when running

--rm to remove container when it stops

```
docker run -d -p 3307:3307 --name mysql --rm mysql
```

```
docker ps
```

Step #4

4. Did it work? Let's interact using cmd-line tool...

```
docker exec -it mysql bash
```

```
mysql -u root -p  
abc123
```

```
show databases;
```

```
exit
```

```
exit
```


```
docker stop mysql
```

```
docker ps
```

Persisting data in the database?

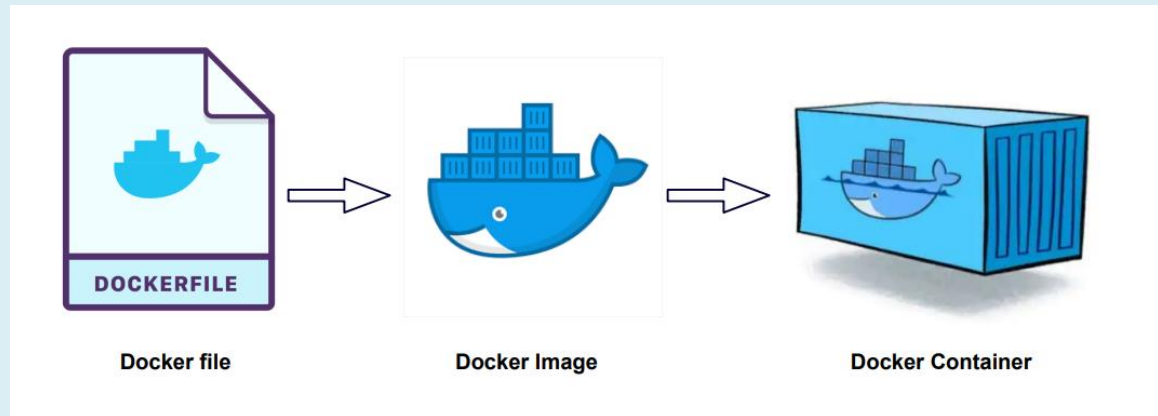
- By default, docker images are stateless --- you do not save data within the docker image
- Instead, you mount a local folder ("volume") and save the data to your local machine with -v option

```
docker run -d -v /data/mysql:/var/lib/mysql ...
```



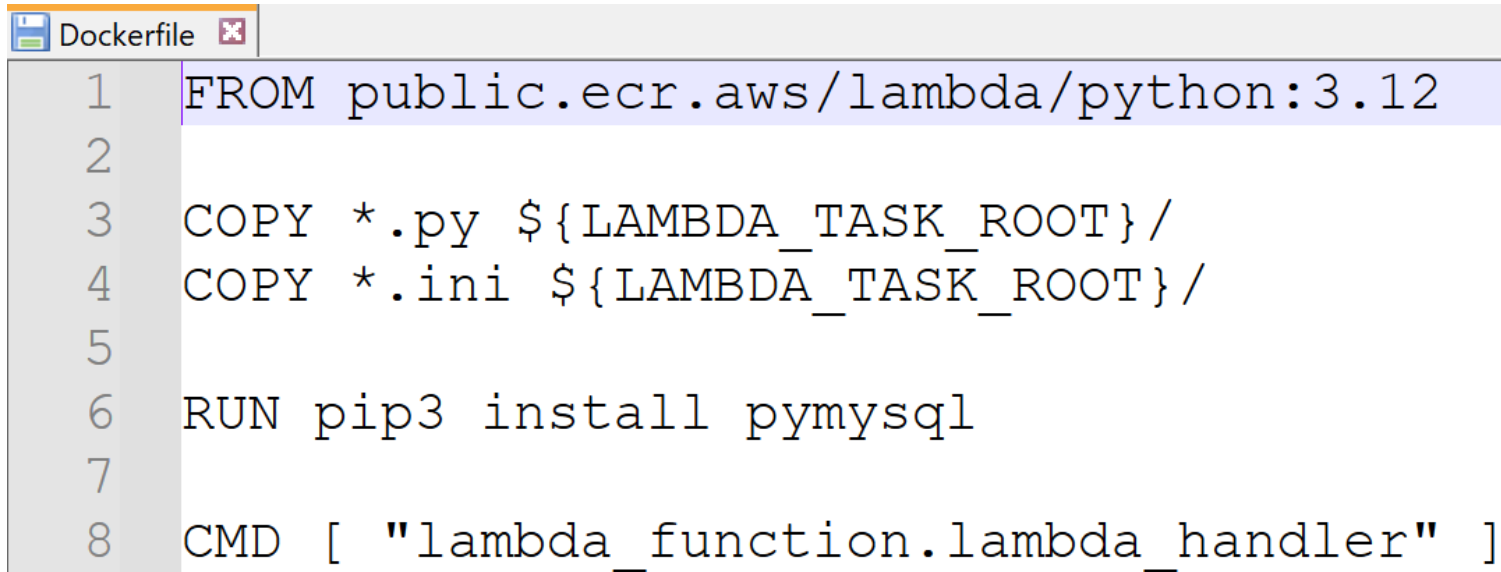
Example #2

- Let's implement a Lambda function in a Docker container to retrieve movies from MovieLens DB running in AWS...



Step #1

1. Create the Dockerfile



```
1 FROM public.ecr.aws/lambda/python:3.12
2
3 COPY *.py ${LAMBDA_TASK_ROOT}/
4 COPY *.ini ${LAMBDA_TASK_ROOT}/
5
6 RUN pip3 install pymysql
7
8 CMD [ "lambda_function.lambda_handler" ]
```

Step #2

2. Build the docker image “lambda-movies”

```
docker build -t lambda-movies .
```

```
docker images
```


Step #3

3. Run the docker image => docker container

-d option is NOT specified so we can see output from lambda

```
docker run -p 8080:8080 --rm lambda-movies
```

Step #4

4. Test locally using POSTMAN desktop app

POST

<http://localhost:8080/2015-03-31/functions/function/invocations>

BODY raw JSON

```
{
  "pathParameters":
  {
    "limit": 5,
    "offset": 50
  }
}
```

```
{
  "statusCode": 200,
  "body": "[[86, \"2006-02-10T00:00:00\", 105, \"de\", 6000000.0, 0.0, \"The Elementary Particles\"], [87, \"1984-05-23T00:00:00\", 118, \"en\", 28000000.0, 333000000.0, \"Indiana Jones and the Temple of Doom\"], [88, \"1987-08-21T00:00:00\", 100, \"en\", 6000000.0, 213954000.0, \"Dirty Dancing\"], [89, \"1989-05-24T00:00:00\", 127, \"en\", 48000000.0, 474172000.0, \"Indiana Jones and the Last Crusade\"], [90, \"1984-11-29T00:00:00\", 105, \"en\", 15000000.0, 316360000.0, \"Beverly Hills Cop\"]]"
}
```

Step #5

5. **Deploy to AWS ECR, create lambda function from ECR image, test within Lambda console...**

Push to ECR, load into Lambda, test

- Switch over to AWS:
 1. *Search for ECR, open ECR console*
 2. *Create a repository with same name as container image*
 3. *Select repository, view "Push command"*
 4. *Open terminal window on your laptop (AWS CLI must be configured - -- we did this in project 01)*
 5. *Execute each of the "push" commands to push image --- skip the command that builds the image, we did that*
 6. *Switch to Lambda console*
 7. *Create new lambda function*
 8. *Select "Create from Image", and select image from ECR drop-down*
 9. *Name function and create*
 10. *Configuration tab: increase timeout*
 11. *Deploy*
 12. *Test --- logs are available from CloudWatch*

That's it, thank you!