

Docker

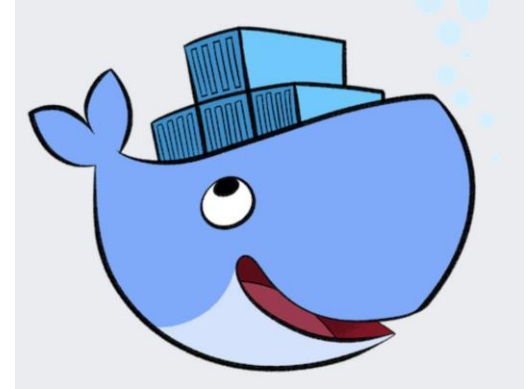
- **What is Docker**
- **Dockerfiles, images, containers**
- **Examples outside AWS**
- **Examples within AWS**



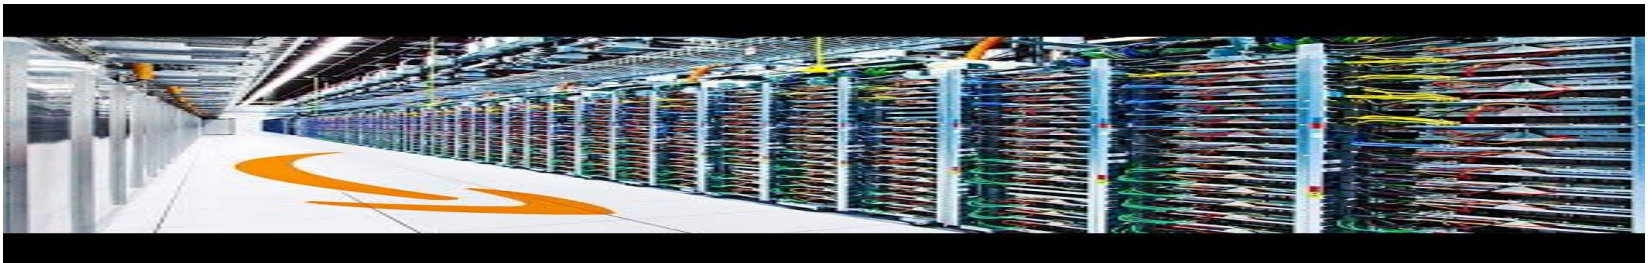
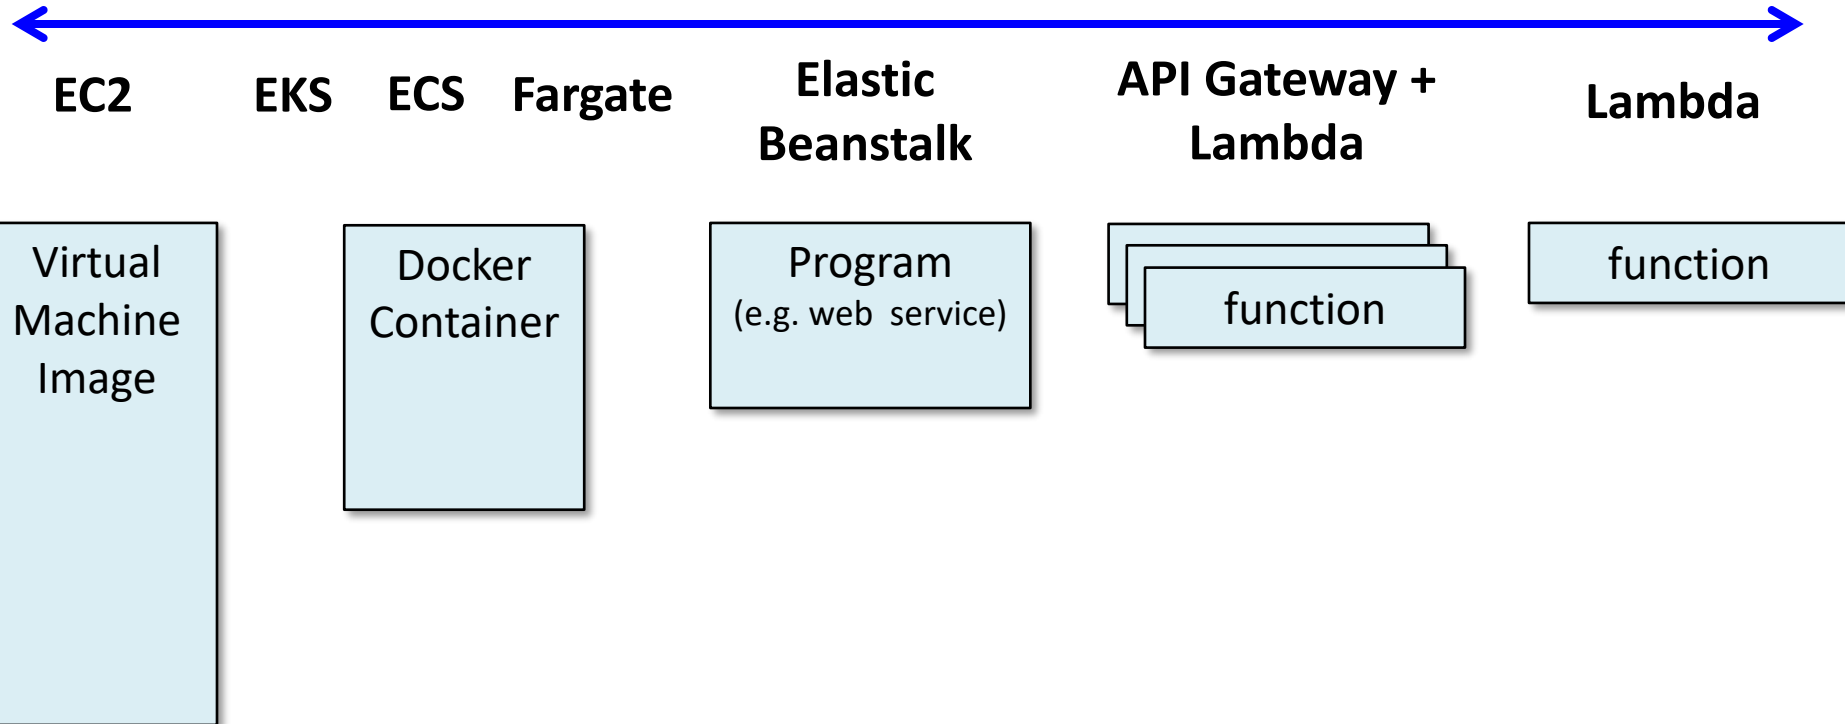
Docker

- **What is Docker?**

- Docker is an open platform for developing, shipping, and running applications
- Docker allows you to run software by installing just two things:
 - Docker desktop
 - Docker image

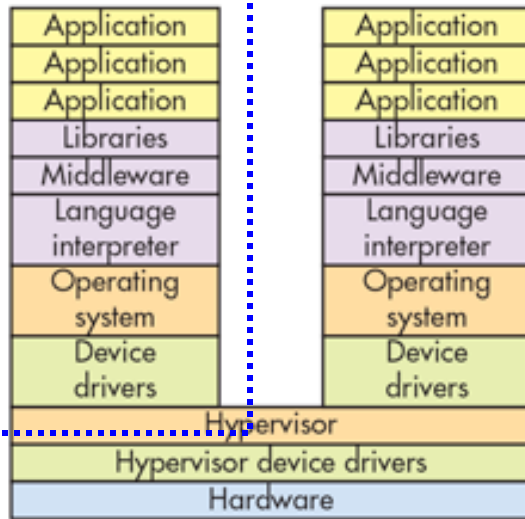


Software packaging options



Software packaging trade-offs

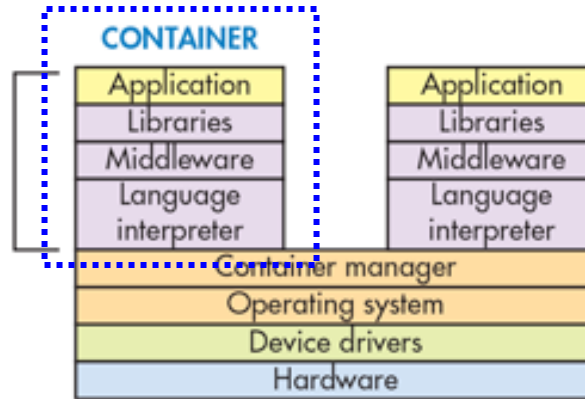
VIRTUAL MACHINE



VIRTUAL MACHINES

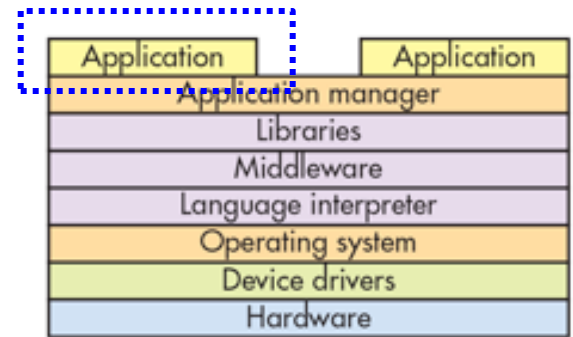
- ✓ Complete control
- ✓ Consistent environment
- ✓ HW-enforced separation/security
- ✗ High space overhead (VM size)
- ✗ Slow cold start since booting OS
(=> slower to scale horizontally)

CONTAINER



CONTAINERS

- ✓ Consistent runtime environment without overhead of a full VM
- ✓ Smaller in size => faster cold start
(=> faster to scale horizontally)
- ✗ Less configurable than a VM
- ✗ Less secure? (SW-enforced)
- ✗ Apps cannot "see" each other, but can "feel" each other (sharing RAM, CPU)



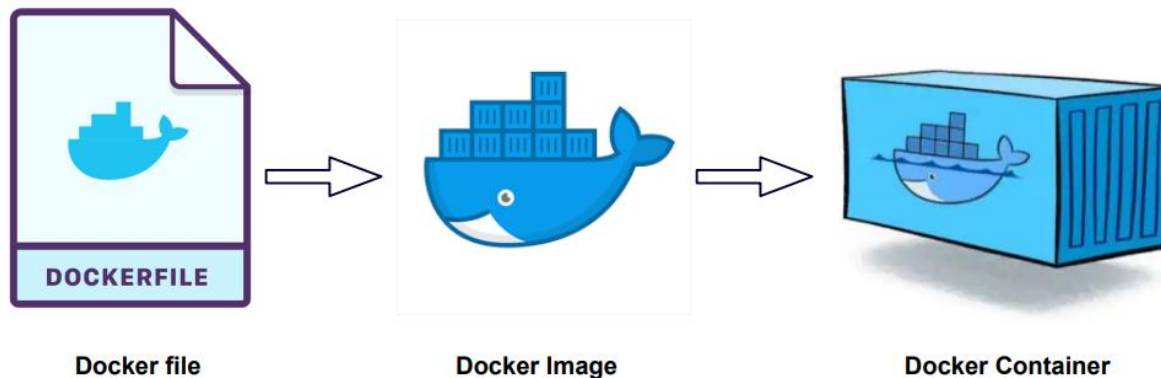
SERVERLESS

- ✓ Easiest to code and deploy
- ✓ Automatic scaling
- ✗ Longer latency (response time)
- ✗ <15 min



Overview

- A **Dockerfile** is a text document that contains the commands needed to build a **docker image**
- Think of a **docker image** as a snapshot of the software – a large file that can be stored and shared
 - *Docker Hub* is the app store for docker images
- A **docker container** runs a docker image



Images vs. Containers

- **Docker desktop** will show your images and containers...

The screenshot shows the Docker Desktop application window. The top bar is blue with the Docker logo, 'docker desktop PERSONAL', a search bar, and several icons. The left sidebar contains a list of navigation items: Containers, Images (selected), Volumes, Builds, Docker Scout, and Extensions. The main area is titled 'Images' with a 'Give feedback' link. Below the title are tabs for 'Local' and 'Hub'. A status bar shows '0 Bytes / 1.03 GB in use' and '6 images'. A search bar and icons for list/table views are present. A table lists the following images:

<input type="checkbox"/>	Name	Tag		Created	Size	Actions	Image ID
<input type="checkbox"/>	project02-server2	latest	○	20 days ago	205.31 MB	▶ ⋮	66e1aef1fc1e
<input type="checkbox"/>	project02-server	latest	○	29 days ago	224.95 MB	▶ ⋮	71a86010375f
<input type="checkbox"/>	web-service-async-demo	latest	○	29 days ago	224.95 MB	▶ ⋮	71a86010375f
<input type="checkbox"/>	project02-client	latest	○	1 month ago	80.3 MB	▶ ⋮	441258a0c512
<input type="checkbox"/>	python3-client	latest	○	1 month ago	80.3 MB	▶ ⋮	441258a0c512
<input type="checkbox"/>	serverless-client	latest	○	1 month ago	80.3 MB	▶ ⋮	441258a0c512

Example #1

- Docker containers can run "interactively", allowing user to run & interact with the software:

```
hummel> ./docker-run.bash
project02-client> python3 main.py
** Welcome to PhotoApp v2 **

What config file to use for this session?
Press ENTER to use default (photoapp-client-config.ini),
otherwise enter name of config file>

>> Enter a command:
  0 => end
  1 => stats
  2 => users
  3 => assets
  4 => download
  5 => download and display
  6 => bucket contents
  7 => add user
  8 => upload
1
bucket status: success
# of users: 5
# of assets: 12

>> Enter a command:
  0 => end
  1 => stats
  2 => users
  3 => assets
  4 => download
  5 => download and display
  6 => bucket contents
  7 => add user
  8 => upload
```

Example #1 files

```
Dockerfile
1 FROM python:3.12.4-alpine3.20
2 #
3 # add bash to alpine Linux:
4 #
5 RUN apk update && apk upgrade
6 RUN apk add --no-cache bash
7 #
8 # turn off history file creation:
9 #
10 RUN echo "export HISTFILE=/dev/null" >> /etc/profile
11 #
12 # add a user (with no pwd) so we don't run as root:
13 #
14 RUN adduser -S user -G users -D
15 #
16 # install any additional python packages we need:
17 #
18 RUN pip3 install requests
19 RUN pip3 install jsons
20
```

```
docker-build.bash
1 #!/bin/bash
2 #
3 # Linux/Mac BASH script to build docker container
4 #
5 docker rmi project02-client
6 docker build -t project02-client .
7
```

```
docker-run.bash
1 #!/bin/bash
2 #
3 # Linux/Mac BASH script to run docker container
4 #
5 docker run -it -u user -w /home/user -v ./home/user --network="host" --rm project02-client bash
6
```


Example #2

- Docker containers can run in the background to offer services, e.g. database server or web server

```
hummel> ./docker-run.bash
22908e6c767fb71fd4b4d9c0299dbf192e3ae4ec224cc1b74e9764ab4d1fb307
hummel>
hummel>
hummel>
hummel> docker ps
CONTAINER ID   IMAGE      COMMAND                  CREATED         STATUS         PORTS
22908e6c767f   docker-mysql  "docker-entrypoint.s..." 4 seconds ago   Up 3 seconds   3306/tcp, 33060/tcp, 0.0.0
.0:3307->3307/tcp   mysql
hummel>
hummel>
hummel>
hummel> docker stop mysql
```

Example #2 files

```
Dockerfile
1 FROM mysql:latest
2 #
3 # set root password for local execution:
4 #
5 ENV MYSQL_ROOT_PASSWORD=abc123
6 #
7 # NOTE: changing to port 3307 since I already have MySQL installed
8 # and running locally on its own. So this docker image is a second
9 # version.
10 #
11 ENV MYSQL_TCP_PORT=3307
12 #
13 # expose the port needed to connect to MySQL server:
14 #
15 EXPOSE 3307
16 #
17 # start server when container runs:
18 #
19 CMD ["mysqld"]
20
```

```
docker-build bash
1 #!/bin/bash
2 #
3 # Linux/Mac BASH script to build docker container
4 #
5 docker rmi docker-mysql
6 docker build -t docker-mysql .
7
```

```
docker-run.bash
1 #!/bin/bash
2 #
3 # Linux/Mac BASH script to run docker container
4 #
5 # NOTE: using port 3307 because the Dockerfile exposes
6 # that port since I already have MySQL running locally.
7 #
8 docker run -d -p 3307:3307 --name mysql --rm docker-mysql
9
```

Docker in AWS

- Docker images are stored in ECR, and then executed via ECS, EKS, Fargate, or Lambda infrastructure



Amazon ECR (Elastic Container Registry)

- ECR is **AWS' Docker container registry service** that allows developers to store, manage and deploy Docker container images.
- **Fully managed and integrated with ECS, EKS, Fargate, Lambda, etc.** making it straightforward to run containerized applications on AWS.



Amazon ECS (Elastic Container Service)

- ECS is a highly scalable, high-performance **container orchestration service**.
- **Used to launch and stop container-based applications** with simple API calls.
- Can be **launched on EC2 instances** (you manage the server) or on **AWS Fargate** (serverless).



Example #3

- Let's run a lambda function using Docker:

```
1 import json
2
3 def lambda_handler(event, context):
4     try:
5         #number1 = int(event['n1'])
6         #number2 = int(event['n2'])
7
8         print("call to add2...")
9
10        params = event["queryStringParameters"]
11
12        number1 = int(params["n1"])
13        number2 = int(params["n2"])
14
15        print("adding", number1, '+', number2)
16
17        sum = number1 + number2
18
19        print("sum:", sum)
20
21        return {
22            'statusCode': 200,
23            'body': json.dumps(sum)
24        }
25
26    except Exception as err:
27        print("***ERROR**")
28        print(str(err))
29
30        return {
31            'statusCode': 500,
32            'body': json.dumps(str(err))
33        }
34
```

Docker files

```
Dockerfile
1 FROM public.ecr.aws/lambda/python:3.12
2
3 COPY lambda_function.py ${LAMBDA_TASK_ROOT}/
4
5 CMD [ "lambda_function.lambda_handler" ]
```

```
docker-build.bash
1 #!/bin/bash
2 #
3 # Linux/Mac BASH script to build docker container
4 #
5 docker rmi lambda-add2
6 docker build -t lambda-add2 .
7
```

Push to ECR, load into Lambda, test

- Switch over to AWS:
 1. *Search for ECR, open ECR console*
 2. *Create a repository with same name as container image*
 3. *Select repository, view "Push command"*
 4. *Open terminal window on your laptop (AWS CLI must be configured - -- we did this in project 01)*
 5. *Execute each of the "push" commands to push image --- skip the command that builds the image, we did that*
 6. *Switch to Lambda console*
 7. *Create new lambda function*
 8. *Select "Create from Image", and select image from ECR drop-down*
 9. *Name function and create*
 10. *Configuration tab: increase timeout*
 11. *Deploy*
 12. *Test --- logs are available from CloudWatch, not Lambda console*

That's it, thank you!